AWS EMR Instructions

These instructions walk you through the process of creating an initial Amazon EMR (Elastic Map Reduce) cluster using **Quick Create** options in the AWS Management Console.

Note, the EMR cluster you set up using these instructions is not meant for a production (secure) environment, and do not cover configuration options in depth. It is meant to help you set up a cluster for class purposes as quickly as possible.

Charges accumulate for cluster you create at the per-second rate for Amazon EMR pricing. The cost will be minimal because the cluster should run for less than a couple of hours after the cluster is provisioned. So it is important that you decommission the cluster as instructed below after you are done with an assignment.

Step 1: Prerequisites

Before you begin setting up your Amazon EMR cluster, make sure that you have completed assignment #1, have an AWS account and understand the basics of working with S3 buckets and associated data objects.

Step 2: AWS Management Console

When you log in to AWS you are presented with the AWS Management Console page. Wherever you are on the site, you can always return to the management console page by clicking on the AWS logo at the top left.

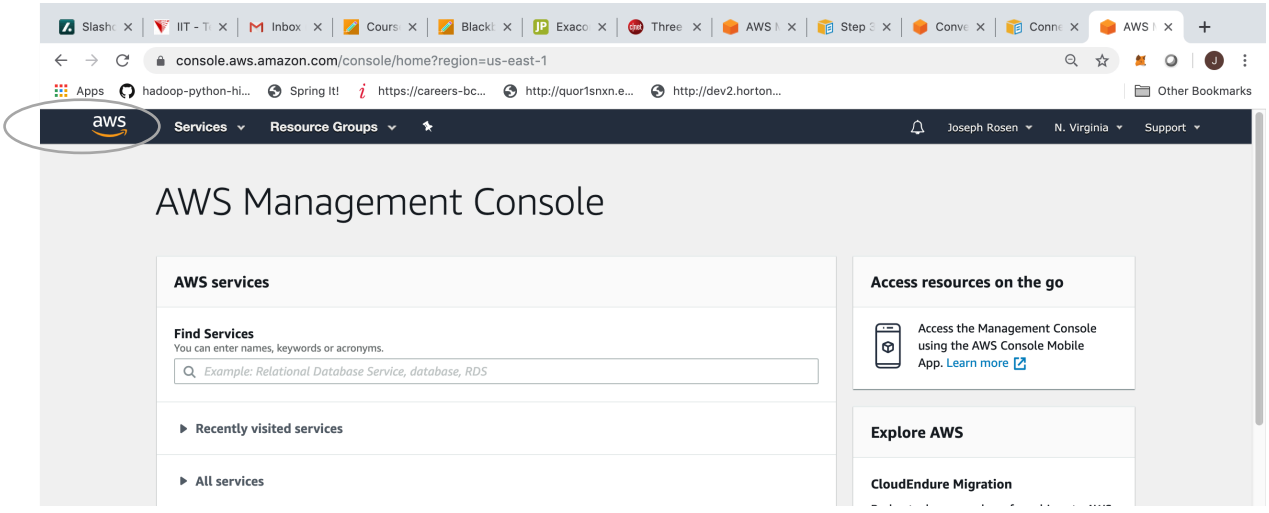

Step 3: Finding Services

We will be making use of several AWS services including

- EC2 – provides computing capability in the form of virtual machines (servers)
- S3 – for object storage
- EMR – Elastic Map Reduce, the Hadoop cluster as a service

When you are on the AWS Management Console page (which we can always get to by clicking the AWS logo), you can find the main page for a service by doing one of the following

1. Type the name of the service whose web page you want to reach into the "Find Services" text box and press Enter/Return
2. If you typed in the name of or used a service recently you might be able to find its name by clicking on "Recently visited services" and then clicking on the name of the desired service
3. If you don't recall the name of the service, then click on "All Services" to get a list and click on the service of interest.
4. Or you can always click on the word "Services" in the upper left of the management console to get a list of services and also type in the one you are looking for.

So in the following steps when you are requested to find some service, you can do the above.

Step 4: Create an Amazon EC2 Key Pair

You must have an Amazon Elastic Compute Cloud (Amazon EC2) key pair to connect to the nodes in your EMR cluster over a secure channel using the Secure Shell (SSH) protocol. We will understand more about SSH below.

1. Find the EC2 service page
2. In the navigation pane, under **NETWORK & SECURITY**, choose **Key Pairs**.

> **Note**
>
> The navigation pane is on the left side of the Amazon EC2 console. If you do not see the pane, it might be minimized; choose the arrow to expand the pane.

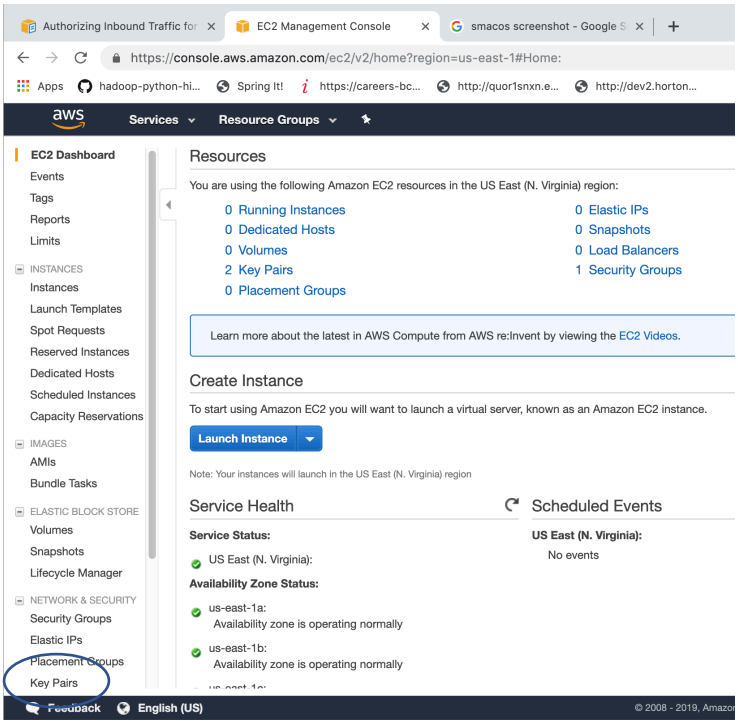

3. Choose **Create Key Pair**.

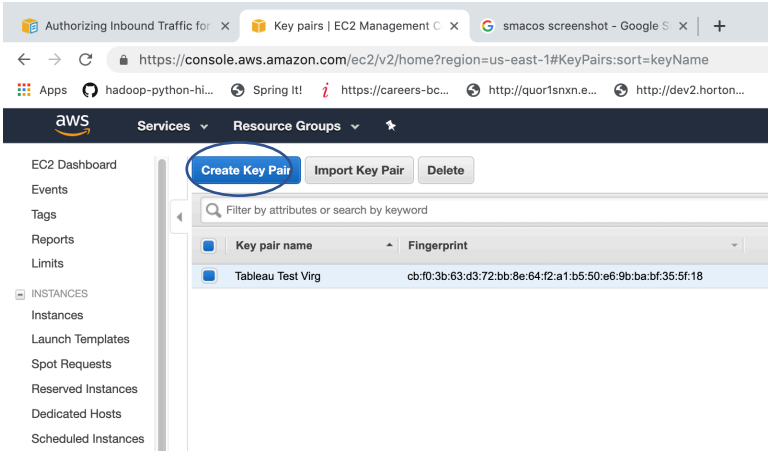

Then you should see a pop up form:

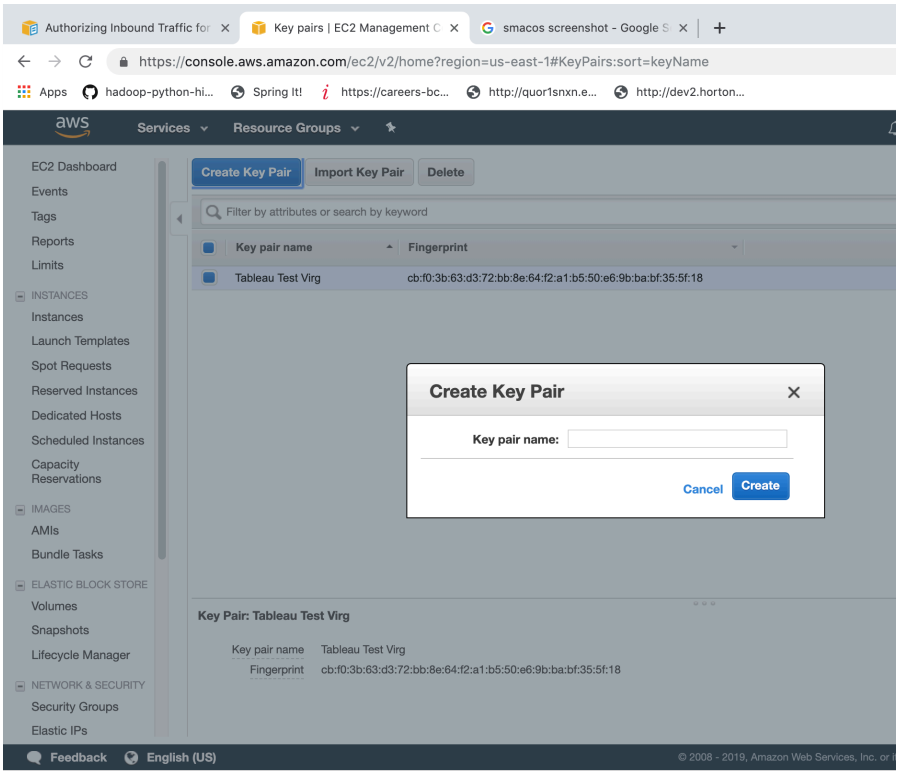

4. For **Key pair name**, enter a name for the new key pair (something like emr-key-pair), and then choose **Create**.

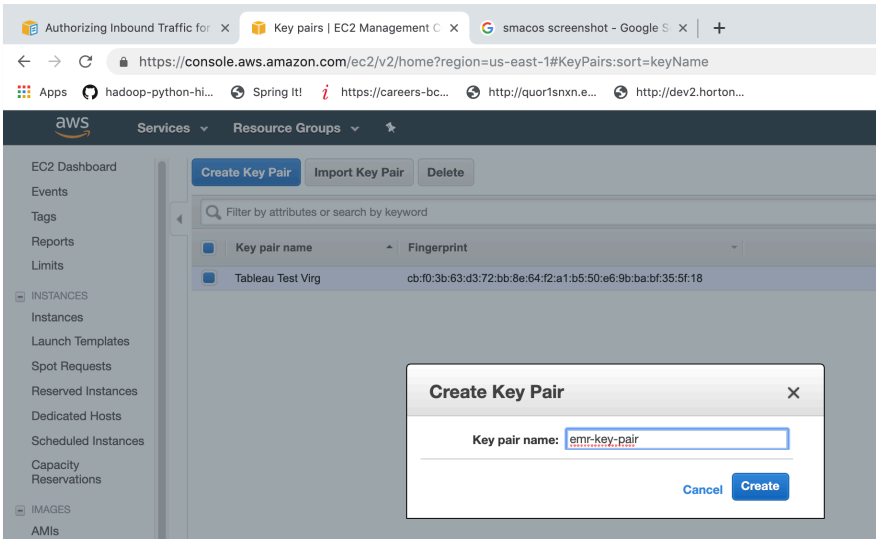

5. The private key file is automatically downloaded by your browser. The base file name is the name you specified as the name of your key pair, and the file name extension is .pem. Save the private key file in a safe place.

   In most cases on the MAC the file will download to the directory /Users/<username>/Downloads

   And on the PC the file will most likely download to

   /c/Users/<username>/Downloads.

   Note, the way I have written the path to the file is formatted for when using the git bash utility.

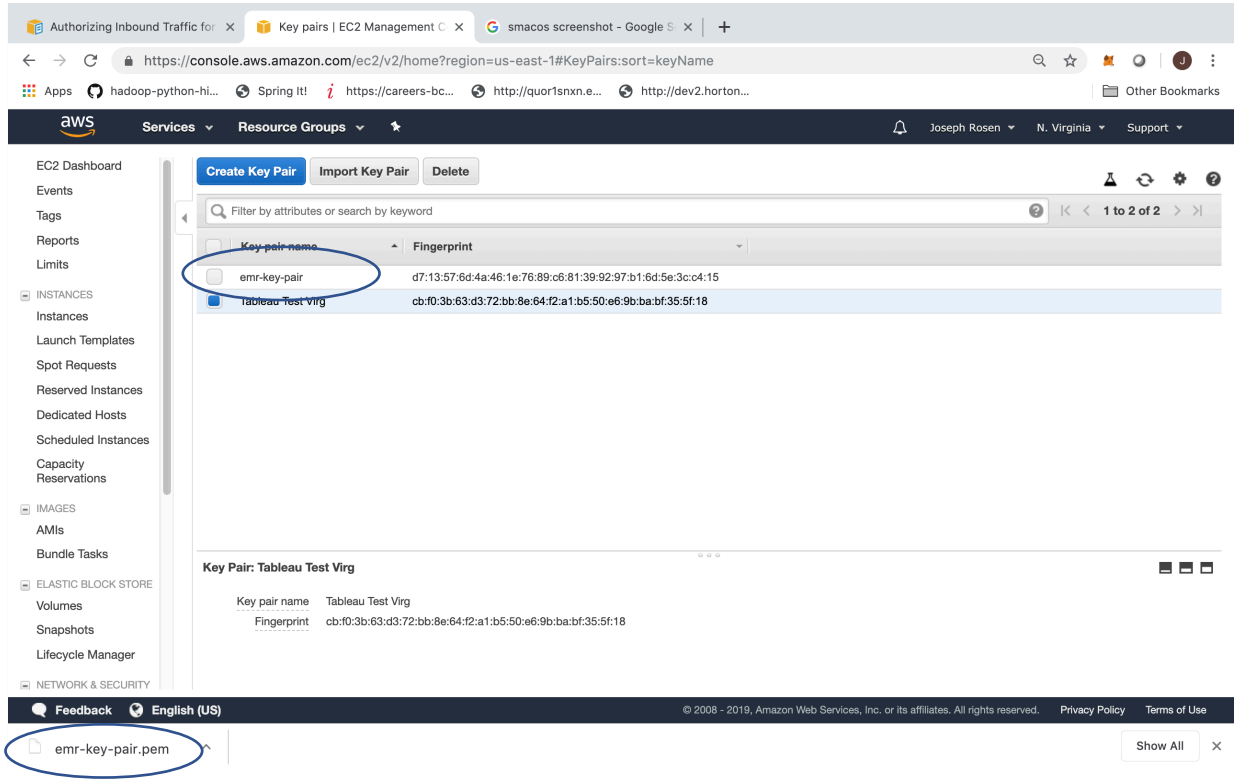

   Important

   This is the only chance for you to save the private key file. You'll need to provide the name of your key pair when you launch an instance and the corresponding private key each time you connect to the instance.

6. So find the directory into which your .pem file has been downloaded and either keep it there or move it to another directory of your choice. You will need to know the path to this file.

7. Using the “terminal” program on the MAC or the “bash” utility on the PC execute the following command to set the permissions of your private key file so that only you can read it. Note, use the appropriate path and file name for your situation.

```
chmod 400 <path-to-file>/emr-key-pair.pem
```

If you do not set these permissions, then you cannot connect to your EMR cluster using this key pair.

Step 5: Launch Your Initial Amazon EMR Cluster

In this step, you launch your initial cluster by using **Quick Options** in the Amazon EMR console and leaving most options to their default values.

To launch the sample Amazon EMR cluster

1. Find the EMR console page
2. Choose **Create cluster**.

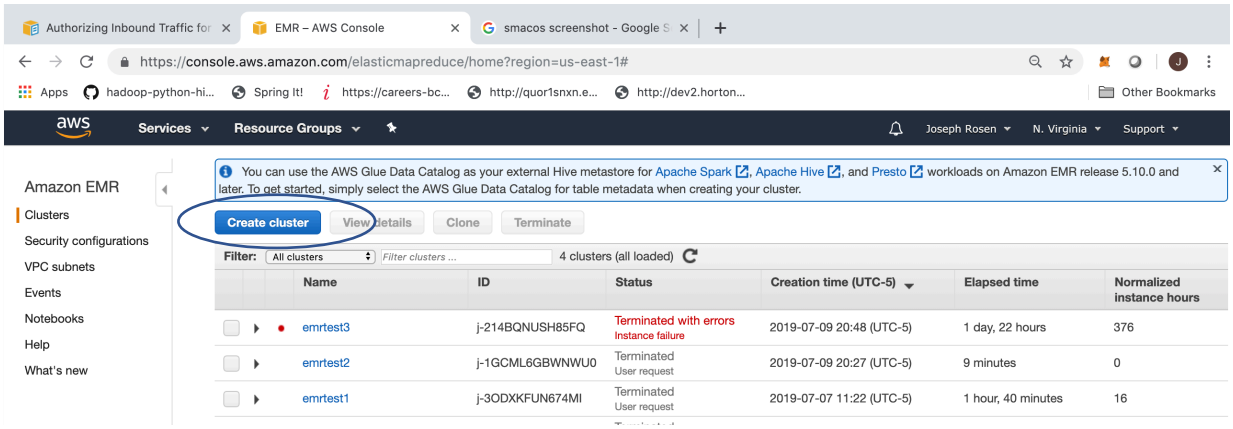

3. On the **Create Cluster - Quick Options** page, accept the default values except for the following fields (see figure on next page):
    - Enter a **Cluster name** that helps you identify the cluster, for example, *My First EMR Cluster*.
    - Under **Hardware configuration** choose:
        - The Instance type as: m4.large
        - The Number of instances as: 2

- Under **Security and access**, choose the **EC2 key pair** that you created in Create an Amazon EC2 Key Pair

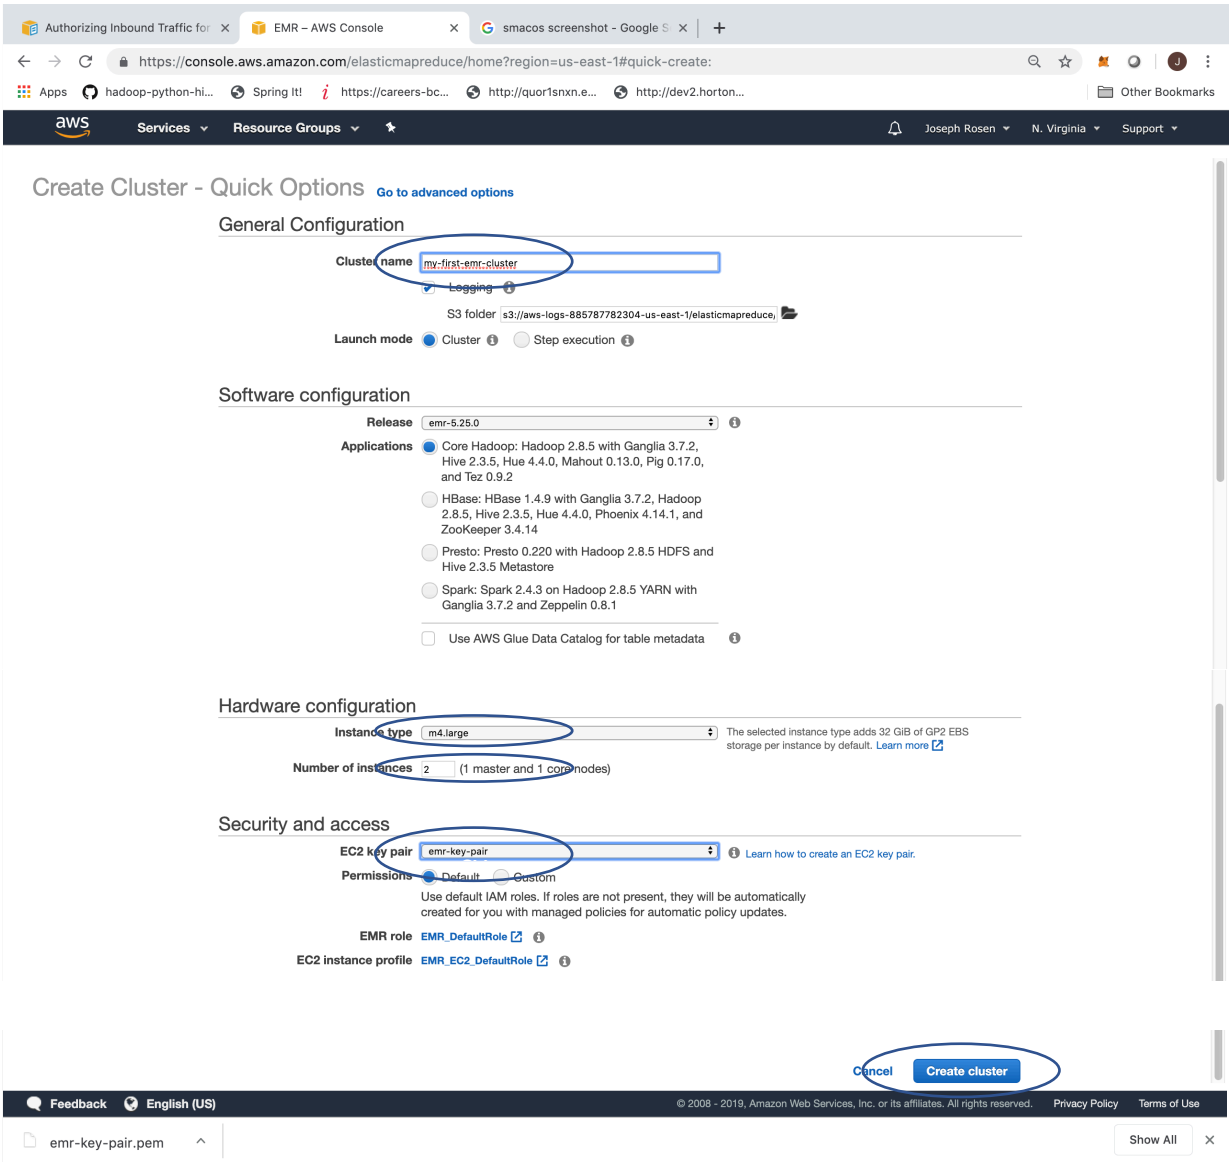

4. Choose **Create cluster**.

The cluster status page with the cluster **Summary** appears (see below). You can use this page to monitor the progress of cluster creation and view details about cluster status. As cluster creation tasks finish, items on the status page update. You may need to choose the refresh icon (circular arrow) on the right or refresh your browser to receive updates.

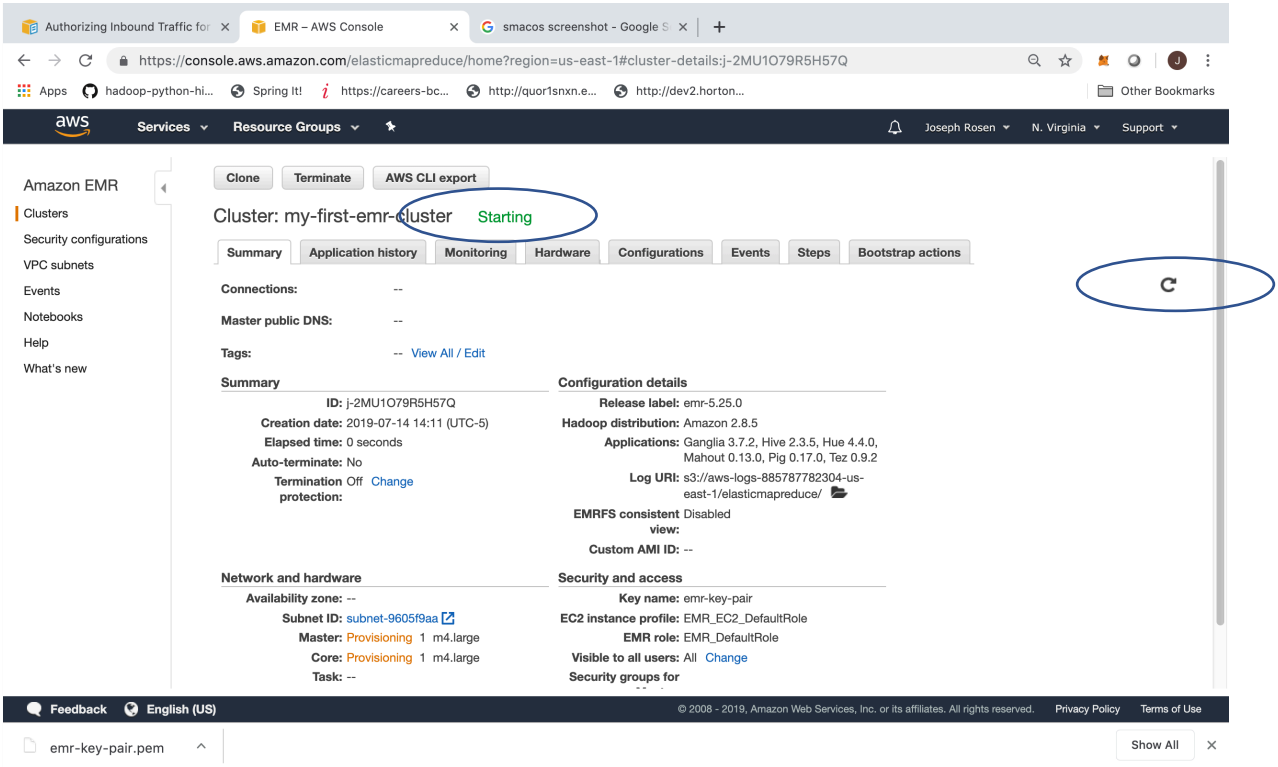

Under **Network and hardware**, find the **Master** and **Core** instance status. The status goes from **Provisioning** to **Bootstrapping** to **Waiting** during the cluster creation process. For more information, see Understanding the Cluster Lifecycle.

As soon as you see the links for **Security groups for Master** and **Security Groups for Core & Task (see below)**, you can move on to the next task, but you may want to wait until the cluster starts successfully and is in the **Waiting** state. The links are blue colored identifiers starting with "sg-" in the Security and Access Area of the page.

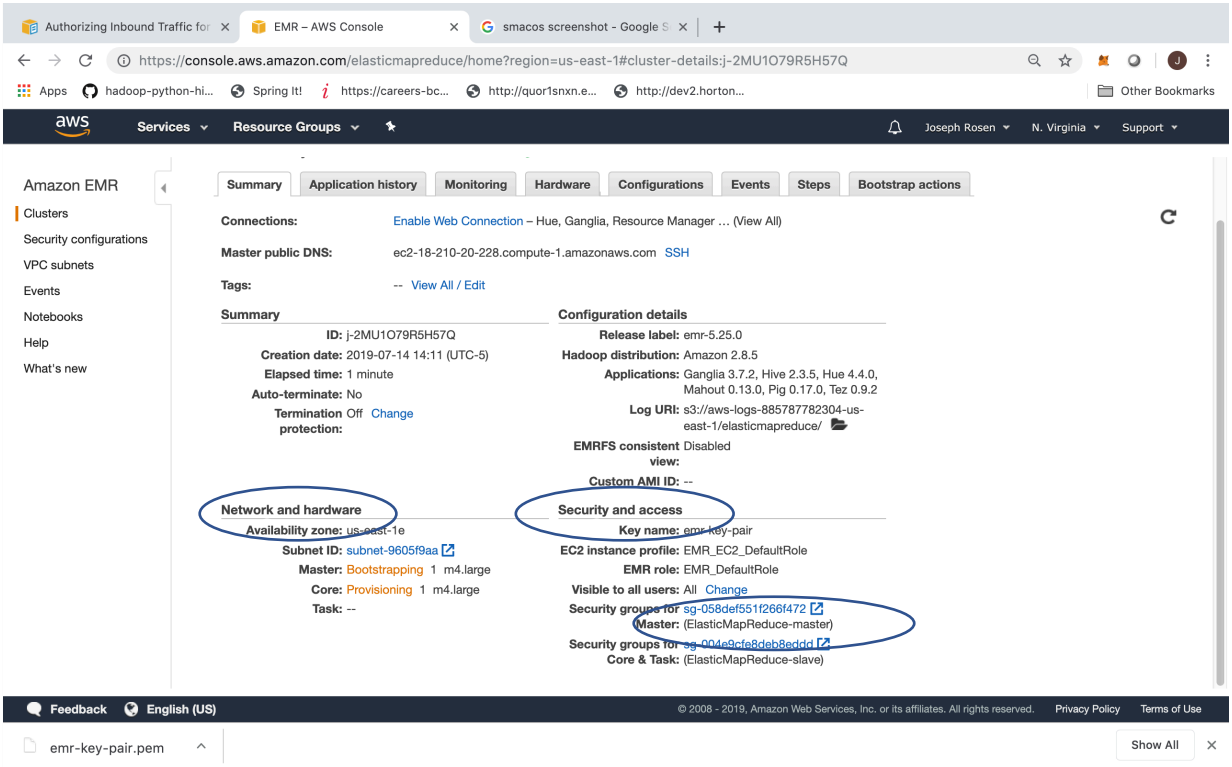

For more information about reading the cluster summary, see View Cluster Status and Details.

Allow SSH Connections to the Cluster From Your Client

Security groups act as virtual firewalls to control inbound and outbound traffic to your cluster. When you create your first cluster, Amazon EMR creates the default Amazon EMR-managed security group associated with the master instance, **ElasticMapReduce-master**, and the security group associated with core and task nodes, **ElasticMapReduce-slave**. To reach the security groups of interest just click on the blue link associated with the Security group for Master entry and you should then see something like the following.

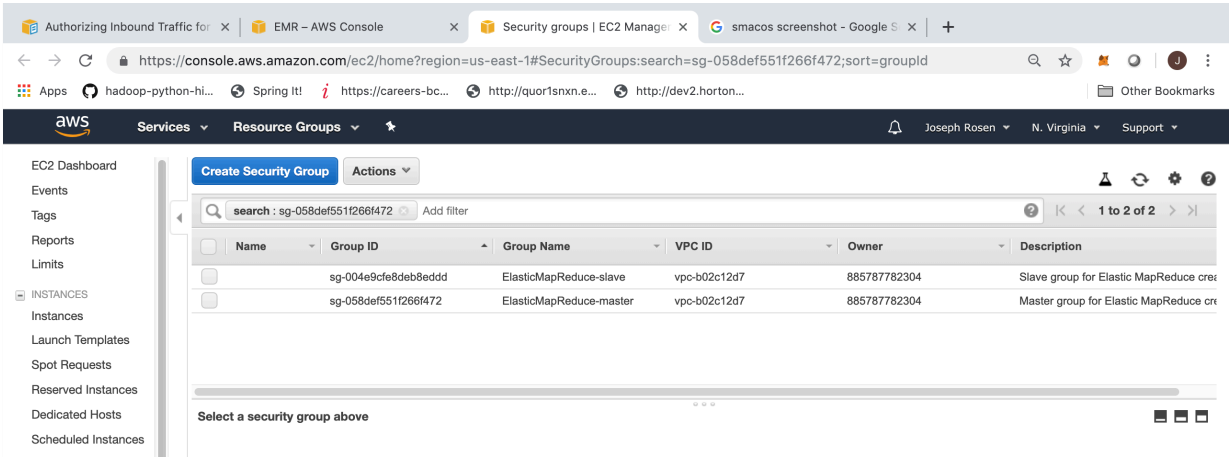

For more information about security groups, see Control Network Traffic with Security Groups and Security Groups for Your VPC in the *Amazon VPC User Guide*.

1. Under **Security and access** choose the **Security groups for Master** link.
2. Choose **ElasticMapReduce-master** from the list.
3. Select the ElasticMapReduce-master by clicking on its row. On the bottom of the screen will appear tabs for this security group. Select the “Inbound” tab.

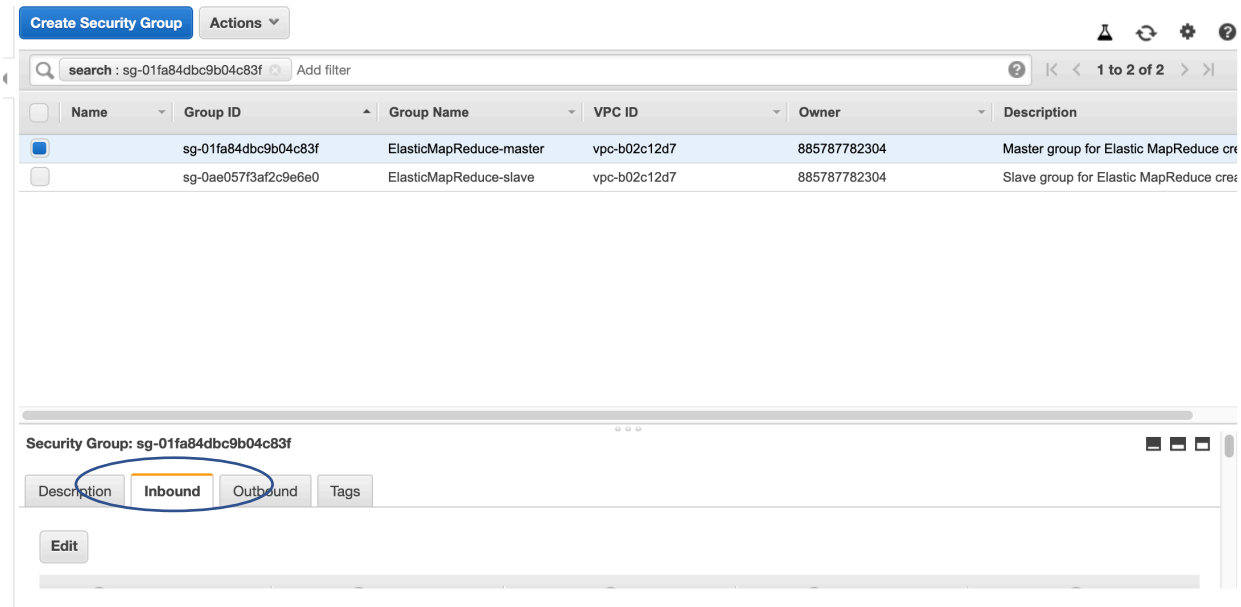

Scroll down to the bottom of the list of inbound rules until you see the “Add Rule” button. Click on it.

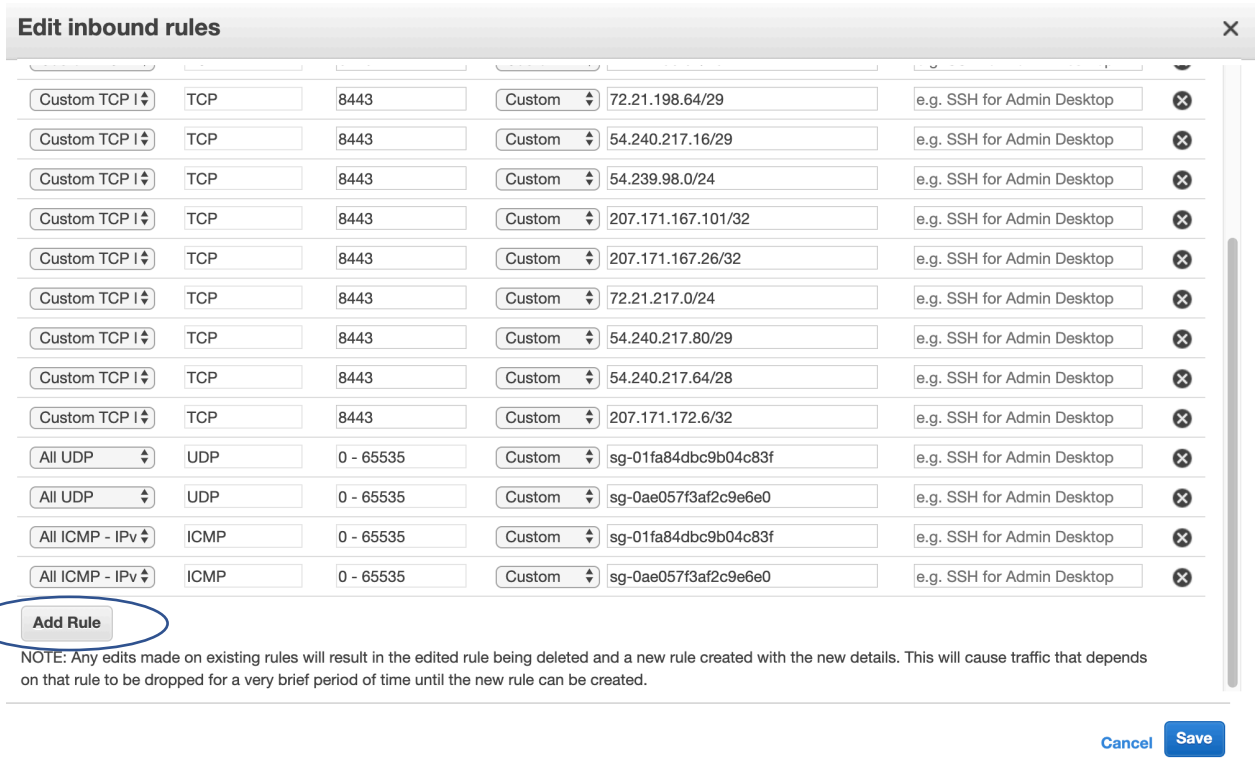

A new row will appear at the bottom of the list which you will modify as follows (also see bwlow):

1. Select the field with label “Custom TCP” which pops up a list of options, select “SSH”. When you do the next field to its left will display the value “TCP” and the next field to the left of that will show “22”.
2. Now select the next field showing the value “Custom” which pops up a list from which you should select “Anywhere”.
3. Now click on the save button.

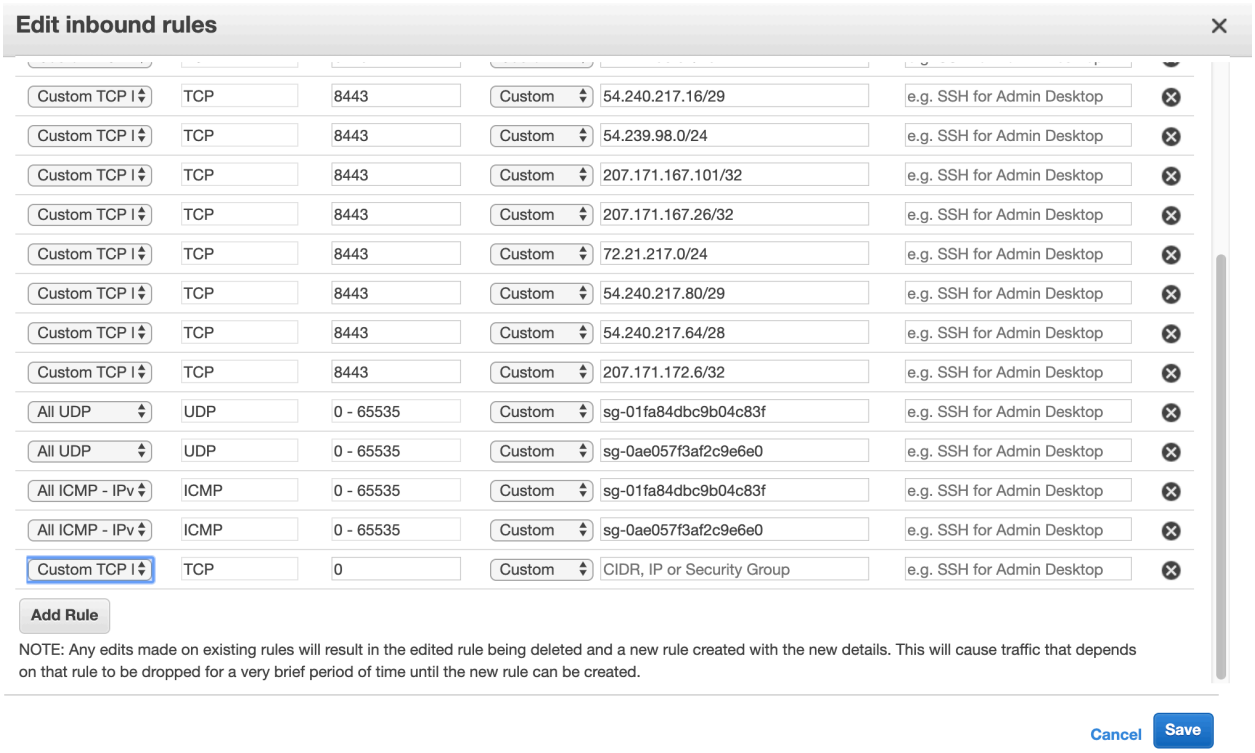
Edit inbound rules
Custom TCP | TCP 8443 Custom 54.240.217.16/29 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.239.98.0/24 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.167.101/32 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.167.26/32 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 72.21.217.0/24 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.240.217.80/29 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.240.217.64/28 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.172.6/32 e.g. SSH for Admin Desktop
All UDP UDP 0 - 65535 Custom sg-01fa84dbc9b04c83f e.g. SSH for Admin Desktop
All UDP UDP 0 - 65535 Custom sg-0ae057f3af2c9e6e0 e.g. SSH for Admin Desktop
All ICMP - IPv ICMP 0 - 65535 Custom sg-01fa84dbc9b04c83f e.g. SSH for Admin Desktop
All ICMP - IPv ICMP 0 - 65535 Custom sg-0ae057f3af2c9e6e0 e.g. SSH for Admin Desktop
Custom TCP | TCP 0 Custom CIDR, IP or Security Group e.g. SSH for Admin Desktop
Add Rule
NOTE: Any edits made on existing rules will result in the edited rule being deleted and a new rule created with the new details. This will cause traffic that depends on that rule to be dropped for a very brief period of time until the new rule can be created.
Cancel
Save

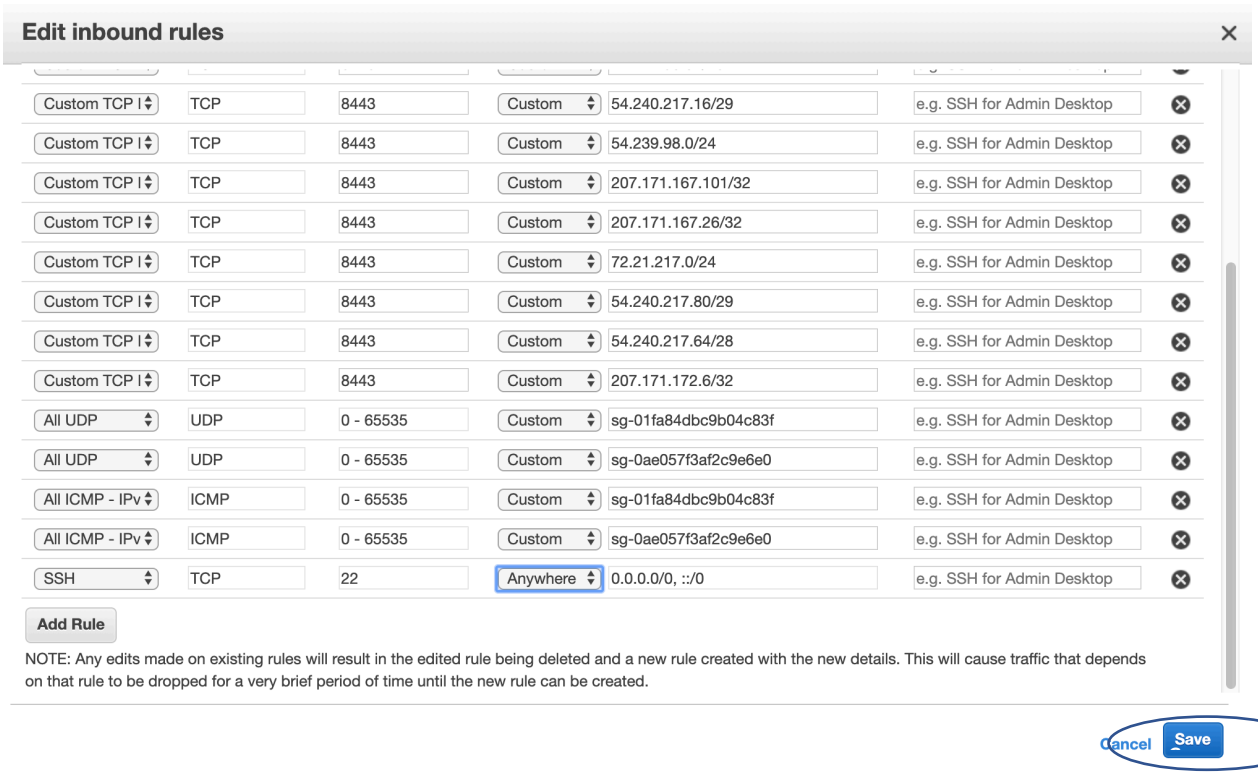
Edit inbound rules
Custom TCP | TCP 8443 Custom 54.240.217.16/29 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.239.98.0/24 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.167.101/32 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.167.26/32 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 72.21.217.0/24 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.240.217.80/29 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 54.240.217.64/28 e.g. SSH for Admin Desktop
Custom TCP | TCP 8443 Custom 207.171.172.6/32 e.g. SSH for Admin Desktop
All UDP UDP 0 - 65535 Custom sg-01fa84dbc9b04c83f e.g. SSH for Admin Desktop
All UDP UDP 0 - 65535 Custom sg-0ae057f3af2c9e6e0 e.g. SSH for Admin Desktop
All ICMP - IPv ICMP 0 - 65535 Custom sg-01fa84dbc9b04c83f e.g. SSH for Admin Desktop
All ICMP - IPv ICMP 0 - 65535 Custom sg-0ae057f3af2c9e6e0 e.g. SSH for Admin Desktop
SSH TCP 22 Anywhere 0.0.0.0/0, ::/0 e.g. SSH for Admin Desktop
Add Rule
NOTE: Any edits made on existing rules will result in the edited rule being deleted and a new rule created with the new details. This will cause traffic that depends on that rule to be dropped for a very brief period of time until the new rule can be created.
Cancel
Save

Step 6: Connect to the Master Node Using SSH

Secure Shell (SSH) is a network protocol you can use to create a secure connection to a remote computer. After you make a connection, the terminal on your local computer behaves as if it is running on the remote computer. Commands you issue locally run on the remote computer, and the command output from the remote computer appears in your terminal window.

When you use SSH with AWS, you are connecting to an EC2 instance, which is a virtual server running in the cloud. When working with Amazon EMR, the most common use of SSH is to connect to the EC2 instance that is acting as the master node of the cluster.

Using SSH to connect to the master node gives you the ability to monitor and interact with the cluster. You can issue Linux commands on the master node, run applications such as Hive and Pig interactively, browse directories, read log files, and so on.

To connect to the master node using SSH, you need the public DNS name of the master node. In addition, the security group associated with the master node must have an inbound rule that allows SSH (TCP port 22) traffic from a source that includes the client where the SSH connection originates (something you did above).

Retrieve the Public DNS Name of the Master Node

You can retrieve the master public DNS name using the Amazon EMR console and the AWS CLI.

To retrieve the public DNS name of the master node using the Amazon EMR console

1. Find the EMR service page.
2. On the **Cluster List** page, select the link for your cluster.
3. Note the **Master public DNS** value that appears at the top of the **Cluster Details** page.

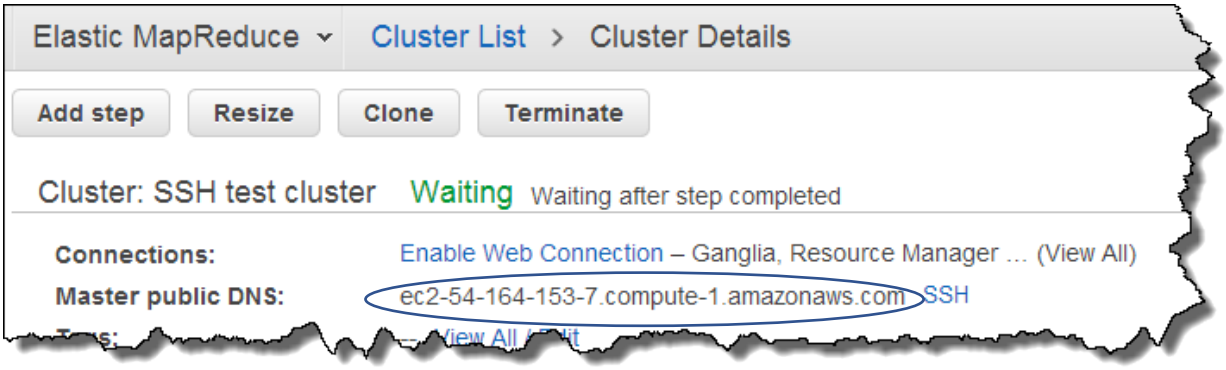

Note

You may also choose the **SSH** link beside the master public DNS name for instructions on creating an SSH connection with the master node, but we will not follow these instructions for the PC

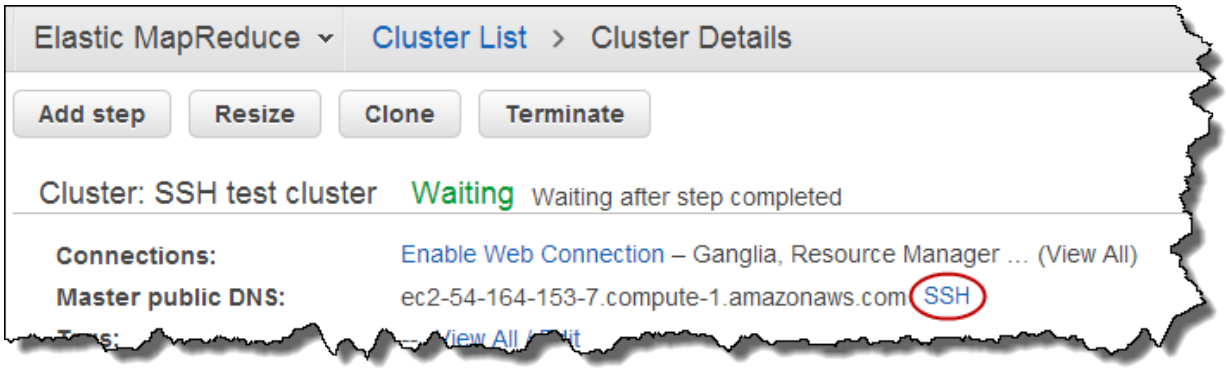

To connect to the Master Node Using SSH and an Amazon EC2 Private Key

Open a terminal window the MAC or use the bash utility on the PC.

1. To establish a connection to the master node, type the following command. Replace *ec2-###-##-##-###.compute-1.amazonaws.com* with the master public DNS name of your cluster and replace *<path-to-file>/mykeypair.pem* with the location and file name of your .pem file.

```
ssh hadoop@ec2-###-##-##-###.compute-1.amazonaws.com -i <path-to-file>/mykeypair.pem
```

   Important

   You must use the login name hadoop when you connect to the Amazon EMR master node; otherwise, you may see an error similar to Server refused our key.

2. A warning states that the authenticity of the host you are connecting to cannot be verified. Type yes to continue.
3. When you are done working on the master node (as you might be at the end of an assignment), type the following command to close the SSH connection.

```
exit
```

Step 7: Terminate the Cluster and Delete the Bucket

After you complete your homework assignment or other project work, you may want to terminate your cluster and delete your Amazon S3 bucket to avoid additional charges.

Terminating your cluster terminates the associated Amazon EC2 instances and stops the accrual of Amazon EMR charges. Amazon EMR preserves metadata information about completed clusters for your reference, at no charge, for two months. The console does not provide a way to delete terminated clusters so that they aren't viewable in the console. Terminated clusters are removed from the cluster when the metadata is removed.

To terminate the cluster

1. Find the EMR service
2. Choose **Clusters**, then choose your cluster.

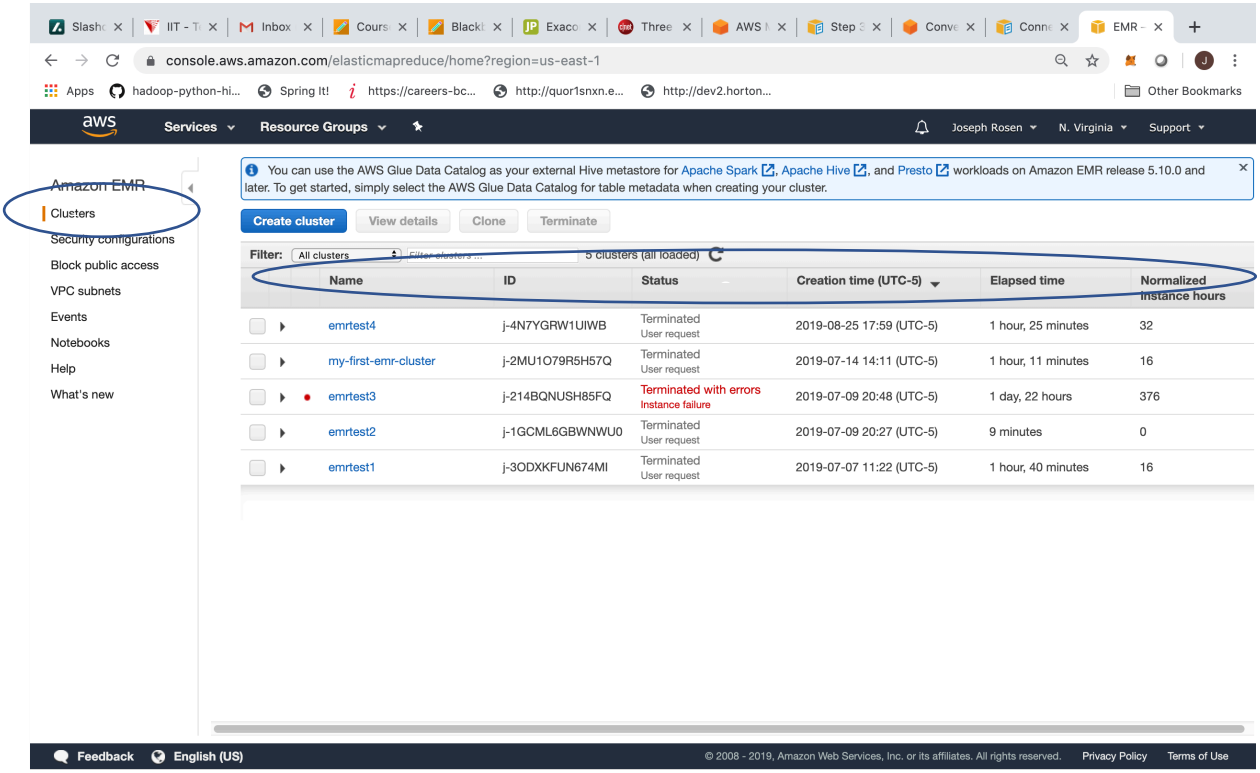

3. Choose Terminate:

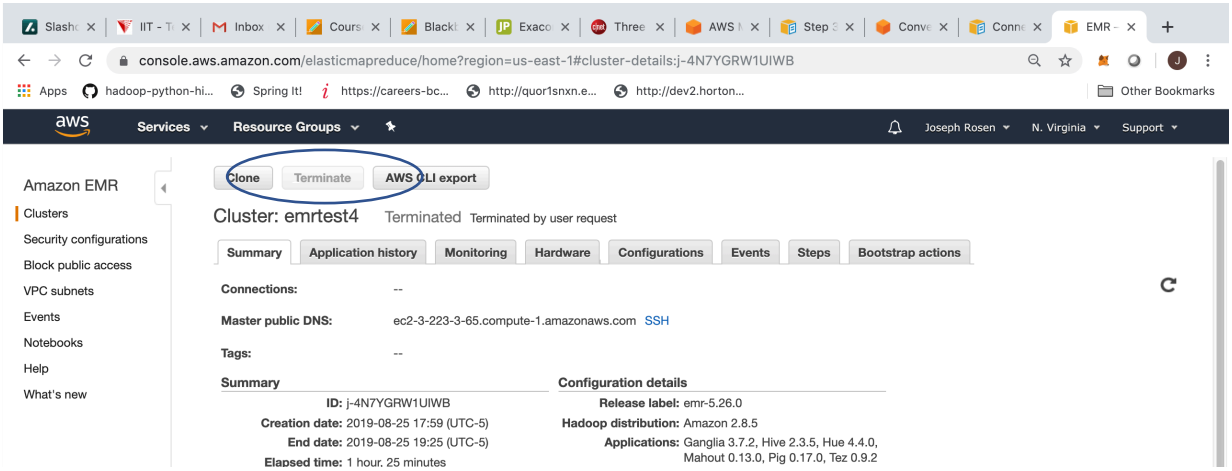

To delete the cluster logging output bucket

1. Find the S3 service
2. Choose the EMR bucket from the list, so that the whole bucket row is selected.
3. Choose delete bucket, type the name of the bucket, and then click **Confirm**.

   For more information about deleting folders and buckets, go to How Do I Delete an S3 Bucket in the *Amazon Simple Storage Service Getting Started Guide*.