

# RegressionSpark

December 10, 2020

```
[7]: !apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://www-us.apache.org/dist/spark/spark-2.4.7/spark-2.4.
→7-bin-hadoop2.7.tgz
!tar xf spark-2.4.7-bin-hadoop2.7.tgz
!pip install -q findspark
```

```
[8]: import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-2.4.7-bin-hadoop2.7"
```

```
[9]: # installing package to plot histogram
import findspark
findspark.init()
from pyspark.sql import SparkSession
spark = SparkSession.builder.master("local[*]").getOrCreate()
```

```
[10]: !pip install pyspark_dist_explore
```

```
Requirement already satisfied: pyspark_dist_explore in /usr/local/lib/python3.6
/dist-packages (0.1.8)
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages
(from pyspark_dist_explore) (1.4.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages
(from pyspark_dist_explore) (1.18.5)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-
packages (from pyspark_dist_explore) (3.2.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-packages
(from pyspark_dist_explore) (1.1.4)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6
/dist-packages (from matplotlib->pyspark_dist_explore) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6
/dist-packages (from matplotlib->pyspark_dist_explore) (1.3.1)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.6/dist-
packages (from matplotlib->pyspark_dist_explore) (0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
/usr/local/lib/python3.6/dist-packages (from matplotlib->pyspark_dist_explore)
(2.4.7)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-
```

packages (from pandas->pyspark\_dist\_explore) (2018.9)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-  
packages (from python-dateutil>=2.1->matplotlib->pyspark\_dist\_explore) (1.15.0)

```
[11]: # Importing the required packages
import pyspark.sql.functions as x
from pyspark.sql.functions import col, count, isnan, when
from pyspark.sql.types import IntegerType
#below code to install files for corelation Matrix
import matplotlib.pyplot as plt
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.stat import Correlation
from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import StandardScaler
from pyspark_dist_explore import hist
```

```
[12]: #importing the Data File into Google Colab
from google.colab import files
files.upload()
```

Output hidden; open in <https://colab.research.google.com> to view.

```
[13]: #Importing the SparkConf, SparkContext
from pyspark import SparkConf, SparkContext
from pyspark.sql import SQLContext
sc = SparkContext.getOrCreate();
sqlContext = SQLContext(sc)
```

```
[14]: #creating a pyspark.sql.dataframe.DataFrame called df by reading the data
      ↳previously imported
df = spark.read.csv('games.csv',inferSchema=True, header =True)
```

```
[15]: # printing the schema of the dataframe
df.printSchema()
```

```
root
 |-- id: integer (nullable = true)
 |-- type: string (nullable = true)
 |-- name: string (nullable = true)
 |-- yearpublished: string (nullable = true)
 |-- minplayers: string (nullable = true)
 |-- maxplayers: string (nullable = true)
 |-- playingtime: integer (nullable = true)
 |-- minplaytime: integer (nullable = true)
 |-- maxplaytime: integer (nullable = true)
 |-- minage: integer (nullable = true)
 |-- usersRated: integer (nullable = true)
```

```

|-- average_rating: double (nullable = true)
|-- bayes_average_rating: double (nullable = true)
|-- total_owners: integer (nullable = true)
|-- total_traders: integer (nullable = true)
|-- total_wanters: integer (nullable = true)
|-- total_wishers: integer (nullable = true)
|-- total_comments: integer (nullable = true)
|-- total_weights: integer (nullable = true)
|-- average_weight: double (nullable = true)

```

```

[16]: # printing the top 10 rows of the dataframe
df.show(10)

```

```

+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
|   id|   type|          name|yearpublished|minplayers|maxplayers|playi
ngtime|minplaytime|maxplaytime|minage|users_rated|average_rating|bayes_average_r
ating|total_owners|total_traders|total_wanters|total_wishers|total_comments|tota
l_weights|average_weight|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+
| 12333|boardgame|    Twilight Struggle|    2005|    2|    2|
180|    180|    180|    13|    20113|    8.33774|
8.22186|    26647|    372|    1219|    5865|    5347|
2562|    3.4785|
|120677|boardgame|    Terra Mystica|    2012|    2|    5|
150|    60|    150|    12|    14383|    8.28798|
8.14232|    16519|    132|    1586|    6277|    2526|
1423|    3.8939|
|102794|boardgame|Caverna: The Cave...|    2013|    1|    7|
210|    30|    210|    12|    9262|    8.28994|
8.06886|    12230|    99|    1476|    5600|    1700|
777|    3.7761|
| 25613|boardgame|Through the Ages:...|    2006|    2|    4|
240|    240|    240|    12|    13294|    8.20407|
8.05804|    14343|    362|    1084|    5075|    3378|
1642|    4.159|
|  3076|boardgame|    Puerto Rico|    2002|    2|    5|
150|    90|    150|    12|    39883|    8.14261|
8.04524|    44362|    795|    861|    5414|    9173|
5213|    3.2943|
| 31260|boardgame|    Agricola|    2007|    1|    5|

```

150	30	150	12	39714	8.11957	
8.03847	47522		837	958	6402	9310
5065	3.616					
124742 boardgame	Android: Netrunner			2012	2	2
45	45	45	14	15281	8.1676	
7.97822	24381		680	627	3244	3202
1260	3.3103					
96848 boardgame	Mage Knight Board...			2011	1	4
150	150	150	14	12697	8.15901	
7.96929	18769		367	1116	5427	2861
1409	4.1292					
84876 boardgame	The Castles of Bu...			2011	2	4
90	30	90	12	15461	8.07879	
7.95011	20558		215	929	3681	3244
1176	3.0442					
72125 boardgame	Eclipse			2011	2	6
200	60	200	14	15709	8.07933	
7.93244	17611		273	1108	5581	3188
1486	3.6359					

```

+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+

```

only showing top 10 rows

```

[17]: # Changing the DataType of columns minplayers,maxplayers,yearpublished from
      →string to Integer Type
df = df.withColumn("minplayers", df["minplayers"].cast(IntegerType()))
df = df.withColumn("maxplayers", df["maxplayers"].cast(IntegerType()))
df = df.withColumn("yearpublished", df["yearpublished"].cast(IntegerType()))

```

```

[18]: # printing the schema of the dataframe
df.printSchema()

```

```

root
|-- id: integer (nullable = true)
|-- type: string (nullable = true)
|-- name: string (nullable = true)
|-- yearpublished: integer (nullable = true)
|-- minplayers: integer (nullable = true)
|-- maxplayers: integer (nullable = true)
|-- playingtime: integer (nullable = true)
|-- minplaytime: integer (nullable = true)
|-- maxplaytime: integer (nullable = true)
|-- minage: integer (nullable = true)
|-- users Rated: integer (nullable = true)
|-- average_rating: double (nullable = true)

```

```

|-- bayes_average_rating: double (nullable = true)
|-- total_owners: integer (nullable = true)
|-- total_traders: integer (nullable = true)
|-- total_wanters: integer (nullable = true)
|-- total_wishers: integer (nullable = true)
|-- total_comments: integer (nullable = true)
|-- total_weights: integer (nullable = true)
|-- average_weight: double (nullable = true)

```

```

[19]: #printing the dimensions of the dataframe
print((df.count(), len(df.columns)))

```

(81312, 20)

```

[20]: #Removing the Rows containing the users Rated below 0
df=df.filter((x.col('users_rated') > 0))

```

```

[21]: #printing the dimensions of the dataframe after the above filtering
print((df.count(), len(df.columns)))

```

(56932, 20)

```

[22]: #removing the following columns "bayes_average_rating", "type", "name", "id"
      →from the dataframe
df=df.select([c for c in df.columns if c not in ["bayes_average_rating",
      →"type", "name", "id"]])

```

```

[23]: #printing the dimensions of the dataframe after the removal of the columns
print((df.count(), len(df.columns)))

```

(56932, 16)

```

[24]: # printing the top 10 rows of the dataframe
df.show(10)

```

```

+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
|yearpublished|minplayers|maxplayers|playingtime|minplaytime|maxplaytime|minage|
users_rated|average_rating|total_owners|total_traders|total_wanters|total_wisher
s|total_comments|total_weights|average_weight|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
|          2005|          2|          2|          180|          180|          180|          13|
20113|          8.33774|          26647|          372|          1219|          5865|

```

5347	2562	3.4785					
	2012	2	5	150	60	150	12
14383	8.28798	16519	132	1586	6277		
2526	1423	3.8939					
	2013	1	7	210	30	210	12
9262	8.28994	12230	99	1476	5600		
1700	777	3.7761					
	2006	2	4	240	240	240	12
13294	8.20407	14343	362	1084	5075		
3378	1642	4.159					
	2002	2	5	150	90	150	12
39883	8.14261	44362	795	861	5414		
9173	5213	3.2943					
	2007	1	5	150	30	150	12
39714	8.11957	47522	837	958	6402		
9310	5065	3.616					
	2012	2	2	45	45	45	14
15281	8.1676	24381	680	627	3244		
3202	1260	3.3103					
	2011	1	4	150	150	150	14
12697	8.15901	18769	367	1116	5427		
2861	1409	4.1292					
	2011	2	4	90	30	90	12
15461	8.07879	20558	215	929	3681		
3244	1176	3.0442					
	2011	2	6	200	60	200	14
15709	8.07933	17611	273	1108	5581		
3188	1486	3.6359					

+-----+-----+-----+-----+-----+-----+-----+

-----+-----+-----+-----+-----+-----+-----

+-----+-----+-----+

only showing top 10 rows

```
[25]: # printing the schema of the dataframe
df.printSchema()
```

```
root
|-- yearpublished: integer (nullable = true)
|-- minplayers: integer (nullable = true)
|-- maxplayers: integer (nullable = true)
|-- playingtime: integer (nullable = true)
|-- minplaytime: integer (nullable = true)
|-- maxplaytime: integer (nullable = true)
|-- minage: integer (nullable = true)
|-- usersRated: integer (nullable = true)
|-- averageRating: double (nullable = true)
```

```

|-- total_owners: integer (nullable = true)
|-- total_traders: integer (nullable = true)
|-- total_wanters: integer (nullable = true)
|-- total_wishers: integer (nullable = true)
|-- total_comments: integer (nullable = true)
|-- total_weights: integer (nullable = true)
|-- average_weight: double (nullable = true)

```

```

[26]: # printing the Count mean stdDev Min Max
df.describe().show()

```

```

+-----+-----+-----+-----+-----+
-- +-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
|summary|      yearpublished|      minplayers|      maxplayers|
playingtime|      minplaytime|      maxplaytime|      minage|
users Rated|      average_rating|      total_owners|      total_traders|
total_wanters|      total_wishers|      total_comments|      total_weights|
average_weight|
+-----+-----+-----+-----+-----+
-- +-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
| count|      56925|      56929|      56929|
56930|      56930|      56930|      56930|      56932|
56932|      56932|      56932|      56932|
56932|      56932|      56932|      56932|
| mean|1874.7589459815547|2.145619982785575|
5.573644364032392|59.935815914280695|57.44503776567715|59.90786931319164|
7.599279817319515|231.21197217733436|6.016053039942434|
374.2626291013841|13.158944003372444|17.99062038923628|
60.53931005409963|70.29883018337667|23.533689313567063|1.2692436924752304|
| stddev| 486.0990365309268|
16.7912994635512|50.757954719008964|406.40176883743374|393.0744116092278| 406.32
4713326961|4.9501535611699605|1363.6812580935118|1.580773749169548|1786.43241221
17926| 46.9707538877096|71.96175435037111|284.07069468630885|338.2672178431332|
138.0081582736383|1.2138650776308242|
| min|      -3500|      0|      0|      0|
0|      0|      0|      0|      1|
1.0|      0|      0|      0|      0|
0|      0|      0.0|
| max|      2017|      2011|      11299|
60120|      60120|      60120|      180|      53680|
25.0|      73188|      1395|      1586|      6402|
11798|      5996|      5.0|

```

```

+-----+-----+-----+-----+-----+-----+
-+-----+-----+-----+-----+-----+-----+
-+-----+-----+-----+-----+-----+-----+
-+-----+-----+-----+-----+-----+-----+

```

```


[29]: #plotting the Histogram
      #!pip install pyspark_dist_explore
      #import pyspark_dist_explore
      fig, ax = plt.subplots()
      #hist(ax,dataset.select("average_rating"), bins = 30, color=['red'])
      hist(ax,df.select("average_rating"), bins = 30, color=['red'])
      #hist(ax, dataset.col withColumn('average_rating')

```

```

[29]: (array([5.9600e+02, 1.0850e+03, 2.1290e+03, 3.8230e+03, 4.2890e+03,
              1.0392e+04, 1.2689e+04, 1.1625e+04, 7.4590e+03, 1.5570e+03,
              8.5600e+02, 4.3100e+02, 0.0000e+00, 0.0000e+00, 0.0000e+00,
              0.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00,
              0.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00,
              0.0000e+00, 0.0000e+00, 0.0000e+00, 0.0000e+00, 1.0000e+00]),
      array([ 1. ,  1.8,  2.6,  3.4,  4.2,  5. ,  5.8,  6.6,  7.4,  8.2,  9. ,
              9.8, 10.6, 11.4, 12.2, 13. , 13.8, 14.6, 15.4, 16.2, 17. , 17.8,
              18.6, 19.4, 20.2, 21. , 21.8, 22.6, 23.4, 24.2, 25. ]),
      <a list of 30 Patch objects>)

```



```

[30]: #Printing the count of NA values present in the DataFrame
      df.select([count(when(col(c).isNull(), c)).alias(c) for c in df.columns]).show()

```



```

+-----+-----+-----+-----+-----+-----+-----+
|yearpublished|minplayers|maxplayers|playingtime|minplaytime|maxplaytime|minage|
users Rated|average_rating|total_owners|total_traders|total_wanters|total_wisher
s|total_comments|total_weights|average_weight|
+-----+-----+-----+-----+-----+-----+-----+
|              7|              3|              3|              2|              2|              2|              2|
0|              0|              0|              0|              0|              0|              0|
0|              0|              0|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+

```

```
[31]: #Dropping the NA values present in the Dataframe
df=df.dropna(how='any')
```

```
[32]: #Printing the count of NA values present in the DataFrame after the removal of NA
df.select([count(when(col(c).isNull(), c)).alias(c) for c in df.columns]).show()
```

```

+-----+-----+-----+-----+-----+-----+-----+
|yearpublished|minplayers|maxplayers|playingtime|minplaytime|maxplaytime|minage|
users Rated|average_rating|total_owners|total_traders|total_wanters|total_wisher
s|total_comments|total_weights|average_weight|
+-----+-----+-----+-----+-----+-----+-----+
|              0|              0|              0|              0|              0|              0|              0|
0|              0|              0|              0|              0|              0|              0|
0|              0|              0|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+

```


```
[33]: #printing the dimensions of the dataframe after the removal of NA
print((df.count(), len(df.columns)))
```

```
(56925, 16)
```

[34]: *#printing the Histogram after NA removal*

```
from pyspark_dist_explore import hist
import matplotlib.pyplot as plt
fig, ax = plt.subplots()
hist(ax, df.select("average_rating"), bins = 30, color=['red'],)
#hist(ax, dataset.col withColumn('average_rating')
```

[34]: (array([ 482., 86., 34., 770., 112., 350., 1565., 379., 880.,  
576., 2760., 1874., 1629., 4107., 2493., 4100., 5666., 3840.,  
4635., 3704., 5517., 3448., 2330., 2693., 845., 652., 726.,  
107., 120., 445.]),  
array([ 1. , 1.3, 1.6, 1.9, 2.2, 2.5, 2.8, 3.1, 3.4, 3.7, 4. ,  
4.3, 4.6, 4.9, 5.2, 5.5, 5.8, 6.1, 6.4, 6.7, 7. , 7.3,  
7.6, 7.9, 8.2, 8.5, 8.8, 9.1, 9.4, 9.7, 10. ]),  
<a list of 30 Patch objects>)



[34]:

[35]: *#Generating the Corelation Matrix*

```
columnsData=df.columns#[c for c in dataset2.columns if c not in_
→['average_rating'] ]
vector_col = "corr_features"
assembler = VectorAssembler(inputCols=columnsData, outputCol=vector_col)
myGraph_vector = assembler.transform(df)#.select(vector_col)
matrix = Correlation.corr(myGraph_vector, vector_col)
matrix1 = Correlation.corr(myGraph_vector, vector_col).collect()[0][0]
```

```
corrmatrix = matrix1.toArray().toList()
print("The generated Correlation Matrix is:: \n {0}".format(corrmatrix))
```

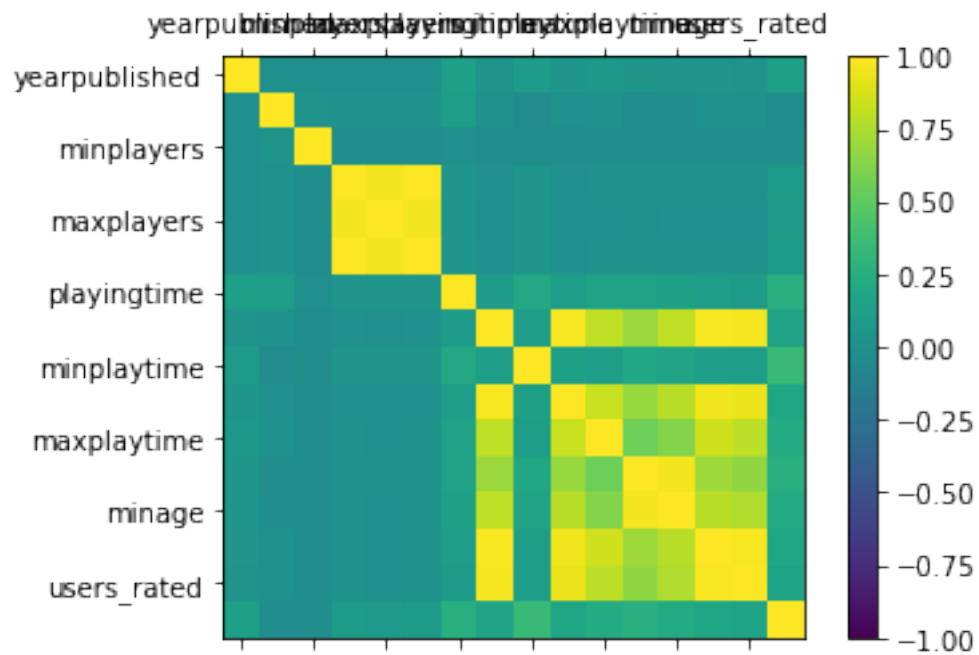
The generated Correlation Matrix is::

```
[[1.0, 0.004372489082671985, 0.0049703856990403325, 0.006232916329865973,
0.005242466181763035, 0.006232916329865973, 0.12522219543828494,
0.037449604506840636, 0.10840383773603023, 0.04773344107604199,
0.06439133752559806, 0.062276921119165635, 0.052525157614782124,
0.043813716831314624, 0.03622198669878126, 0.12530663336643624],
[0.004372489082671985, 1.0, 0.03718879898920534, 0.02428356613412483,
0.0265323496473018, 0.02428356613412483, 0.11117982079071105,
0.020313525169298213, -0.032858165438786856, 0.014536490514991217,
0.024624729226298204, -0.008714669167627476, -0.0038365658456161987,
0.02297155625899625, 0.019011603278792415, -0.022024292273003405],
[0.0049703856990403325, 0.03718879898920534, 1.0, -0.0010206921408426179,
-0.0009504277066825062, -0.0010206921408426179, 0.004521518981639495,
-0.0008609769588981912, -0.008339907100546615, -0.0009813364242238398,
-0.002386341572889714, -0.0035900367546440105, -0.00250317090907312,
-0.0015057172595959244, -0.0014607609885830214, -0.013453775274283387],
[0.006232916329865973, 0.02428356613412483, -0.0010206921408426179, 1.0,
0.9679090330315554, 0.9999999999999998, 0.053442139957624514,
0.010913960113744634, 0.04898751699418791, 0.014539484718191427,
0.01927338845715717, 0.024295985122396398, 0.020303339878855887,
0.016896922591798955, 0.0177950776252157, 0.09090481965245623],
[0.005242466181763035, 0.0265323496473018, -0.0009504277066825062,
0.9679090330315554, 1.0, 0.9679090330315554, 0.05242629945411761,
0.005508101217625019, 0.04397794536793453, 0.008981488408351649,
0.014492718173379271, 0.01749777640721583, 0.012743733653785828,
0.010052547621304374, 0.010177700977191314, 0.0844154545516134],
[0.006232916329865973, 0.02428356613412483, -0.0010206921408426179,
0.9999999999999998, 0.9679090330315554, 1.0, 0.053442139957624514,
0.010913960113744634, 0.04898751699418791, 0.014539484718191427,
0.01927338845715717, 0.024295985122396398, 0.020303339878855887,
0.016896922591798955, 0.0177950776252157, 0.09090481965245623],
[0.12522219543828494, 0.11117982079071105, 0.004521518981639495,
0.053442139957624514, 0.05242629945411761, 0.053442139957624514, 1.0,
0.0978216192262467, 0.20988344866685532, 0.11848439964079609,
0.15220440418785858, 0.15399382162327566, 0.13739024721262805,
0.11331719837221133, 0.09701953445520294, 0.25948447050554835],
[0.037449604506840636, 0.020313525169298213, -0.0008609769588981912,
0.010913960113744634, 0.005508101217625019, 0.010913960113744634,
0.0978216192262467, 1.0, 0.11257620249092382, 0.9776642994358582,
0.8078747615566414, 0.7024387933993025, 0.8051192766502231, 0.9787082100014113,
0.975723832105259, 0.15061170841954058], [0.10840383773603023,
-0.032858165438786856, -0.008339907100546615, 0.04898751699418791,
0.04397794536793453, 0.04898751699418791, 0.20988344866685532,
0.11257620249092382, 1.0, 0.13749262146291352, 0.11947807651217442,
```

0.19657551824080938, 0.1713853185458103, 0.12372990572424777,  
0.10970326962727345, 0.35111626853661754], [0.04773344107604199,  
0.014536490514991217, -0.0009813364242238398, 0.014539484718191427,  
0.008981488408351649, 0.014539484718191427, 0.11848439964079609,  
0.9776642994358582, 0.13749262146291352, 1.0, 0.829879539373002,  
0.6883883341927931, 0.7888253085245929, 0.9540266393812632, 0.9384231749868207,  
0.18336003786744906], [0.06439133752559806, 0.024624729226298204,  
-0.002386341572889714, 0.01927338845715717, 0.014492718173379271,  
0.01927338845715717, 0.15220440418785858, 0.8078747615566414,  
0.11947807651217442, 0.829879539373002, 1.0, 0.555853036766878,  
0.6314542040905395, 0.8565287058577824, 0.8015616561094298,  
0.22733183308801955], [0.062276921119165635, -0.008714669167627476,  
-0.0035900367546440105, 0.024295985122396398, 0.01749777640721583,  
0.024295985122396398, 0.15399382162327566, 0.7024387933993025,  
0.19657551824080938, 0.6883883341927931, 0.555853036766878, 1.0,  
0.9610027518582096, 0.7096534539020688, 0.6676146441610739, 0.2542429847235272],  
[0.052525157614782124, -0.0038365658456161987, -0.00250317090907312,  
0.020303339878855887, 0.012743733653785828, 0.020303339878855887,  
0.13739024721262805, 0.8051192766502231, 0.1713853185458103, 0.7888253085245929,  
0.6314542040905395, 0.9610027518582096, 1.0, 0.7885428287299897,  
0.7646815806776168, 0.21903535783186218], [0.043813716831314624,  
0.02297155625899625, -0.0015057172595959244, 0.016896922591798955,  
0.010052547621304374, 0.016896922591798955, 0.11331719837221133,  
0.9787082100014113, 0.12372990572424777, 0.9540266393812632, 0.8565287058577824,  
0.7096534539020688, 0.7885428287299897, 1.0, 0.9782125987665806,  
0.1809582653627293], [0.03622198669878126, 0.019011603278792415,  
-0.0014607609885830214, 0.0177950776252157, 0.010177700977191314,  
0.0177950776252157, 0.09701953445520294, 0.975723832105259, 0.10970326962727345,  
0.9384231749868207, 0.8015616561094298, 0.6676146441610739, 0.7646815806776168,  
0.9782125987665806, 1.0, 0.16101794935654037], [0.12530663336643624,  
-0.022024292273003405, -0.013453775274283387, 0.09090481965245623,  
0.0844154545516134, 0.09090481965245623, 0.25948447050554835,  
0.15061170841954058, 0.35111626853661754, 0.18336003786744906,  
0.22733183308801955, 0.2542429847235272, 0.21903535783186218,  
0.1809582653627293, 0.16101794935654037, 1.0]]


[36]:

```
# Generating the HeatMap of the Data
figure=plt.figure(1)
figure1=figure.add_subplot(111)
figure1.set_xticklabels(['']+columnsData)
figure1.set_yticklabels(['']+columnsData)
figure.colorbar(figure1.matshow(corrmatrix,vmax=1,vmin=-1))
plt.show()
```

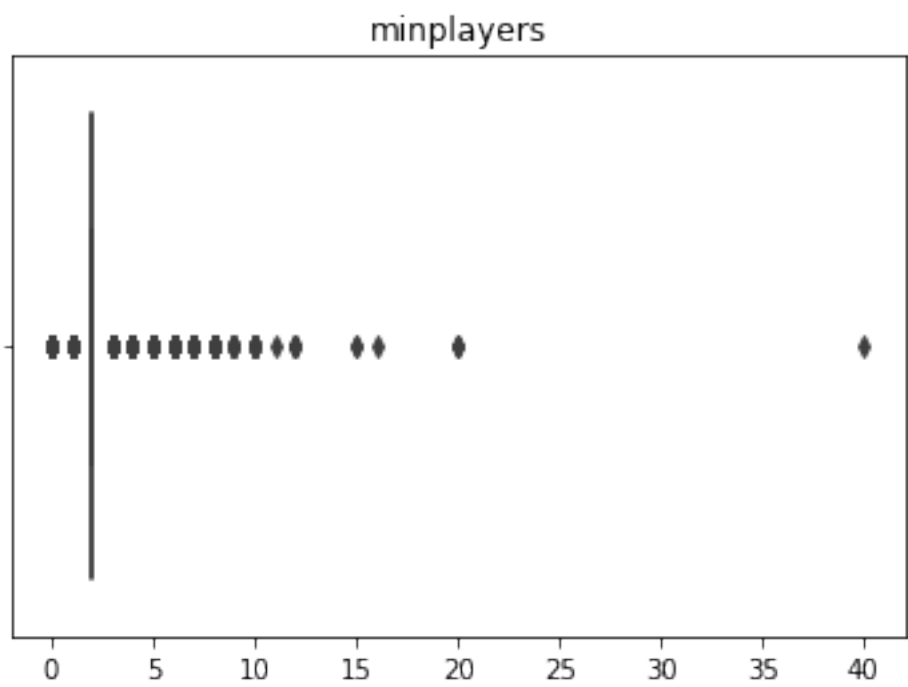
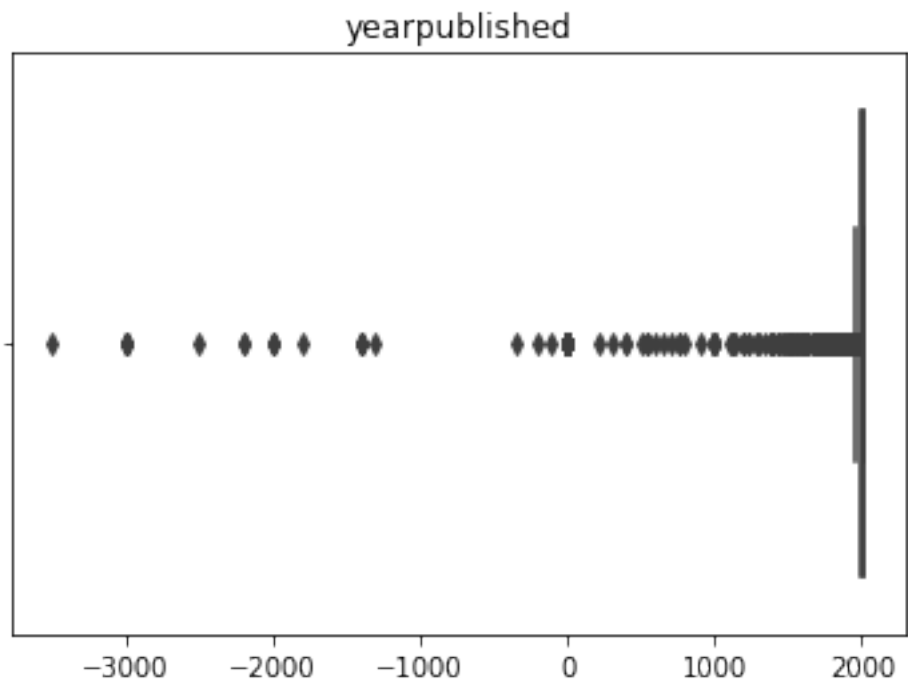



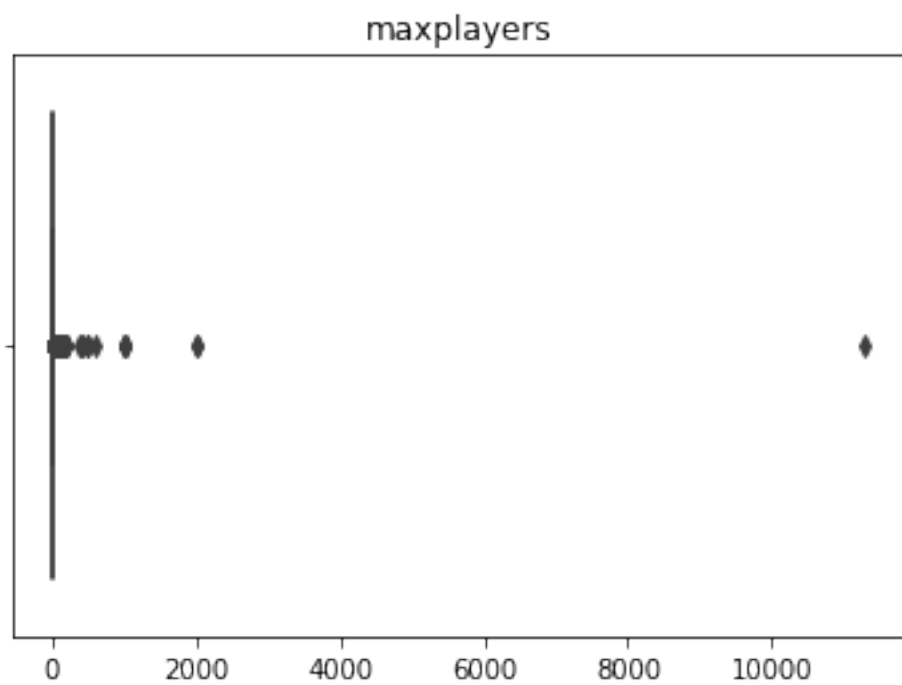
```
[37]: # Plotting the pairplot
import seaborn as sns
sns.pairplot(df.toPandas())
```

```
[37]: <seaborn.axisgrid.PairGrid at 0x7f3f5ca5c630>
```


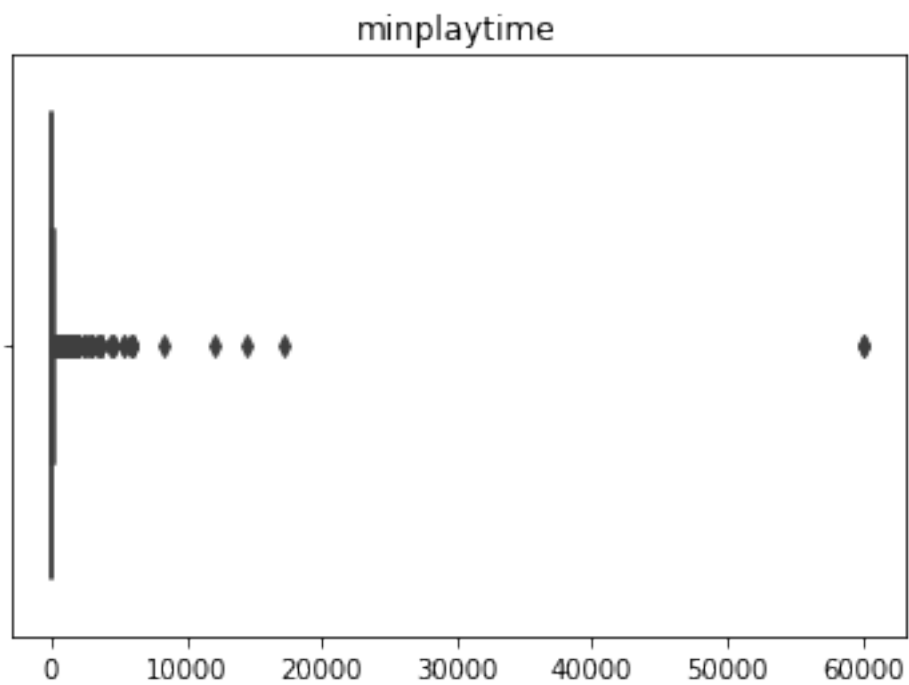




```
[38]: #Plotting the Box plots
for i in columnsData:
    l = df.select(i).collect()
    l = [r[0] for r in l]
    plt.figure()
    b = sns.boxplot(x = l)
    plt.title(i)
    plt.show()
```

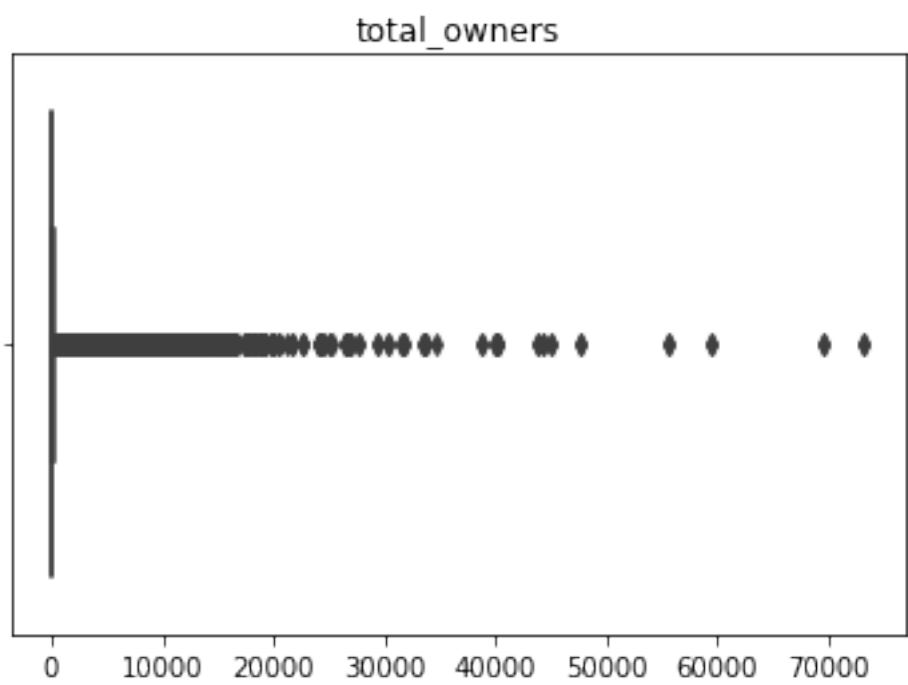
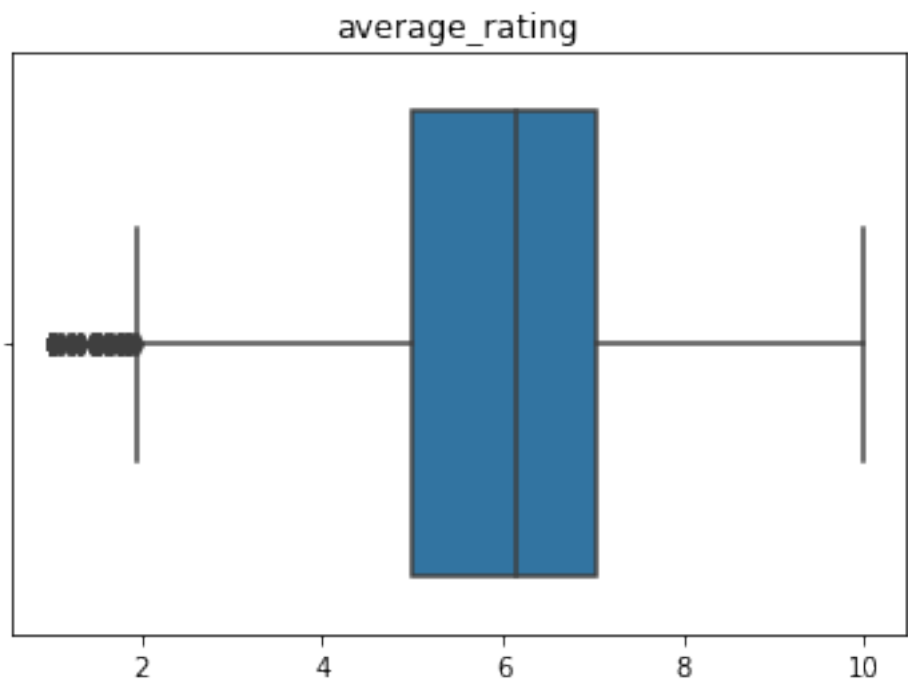


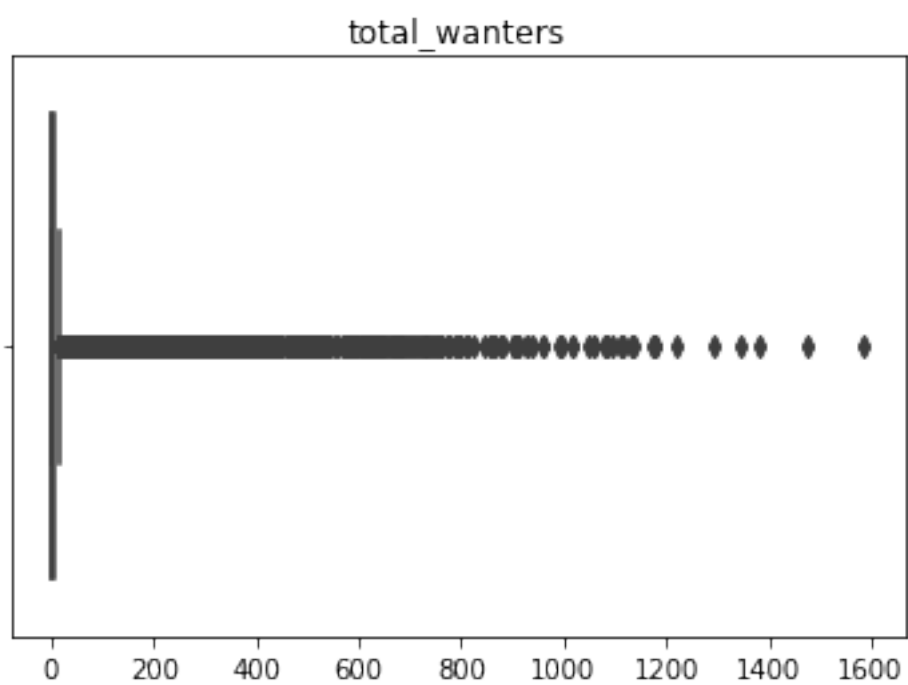
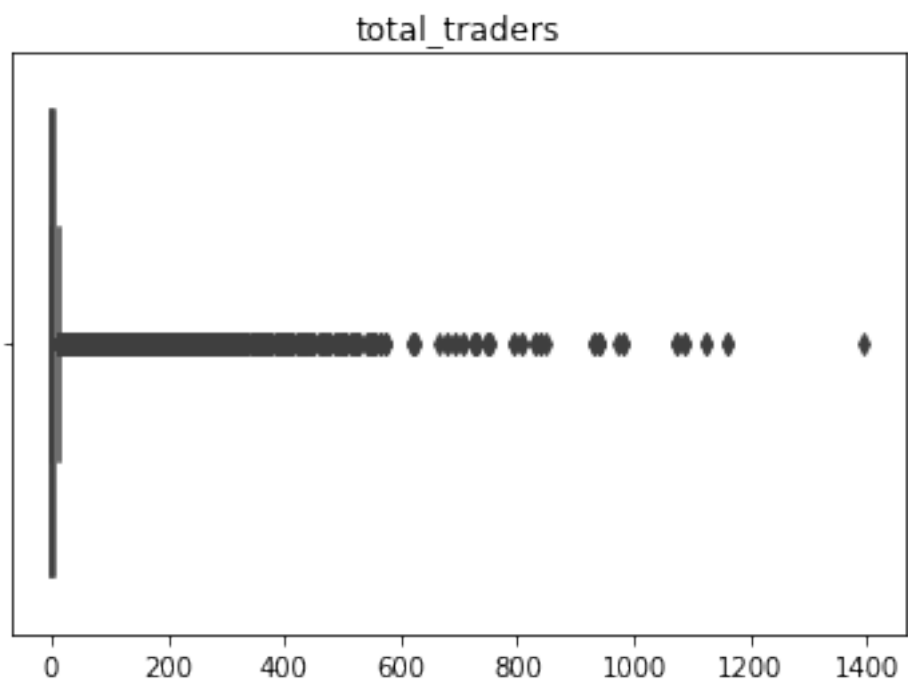



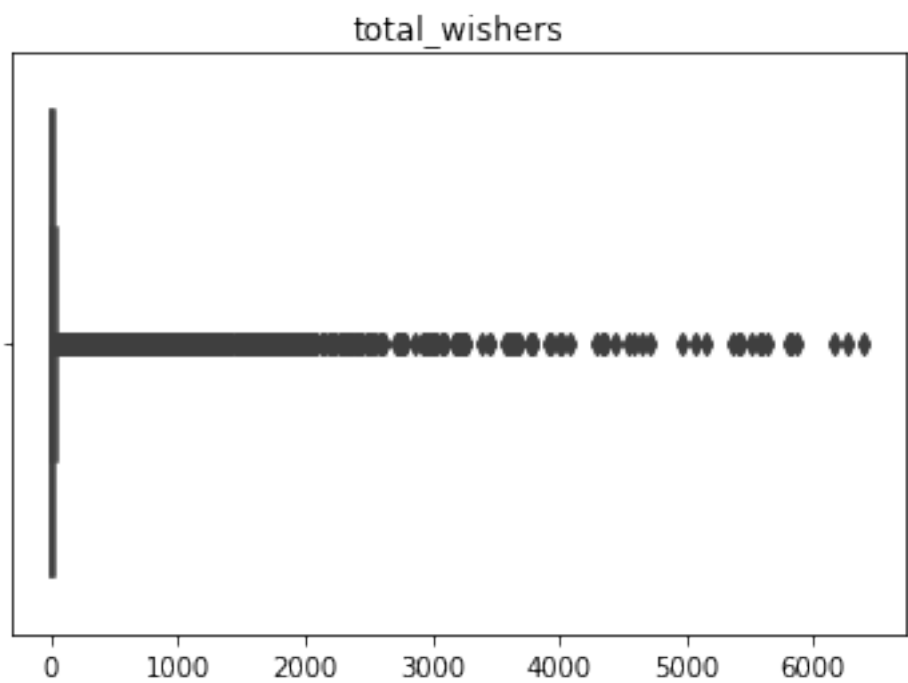


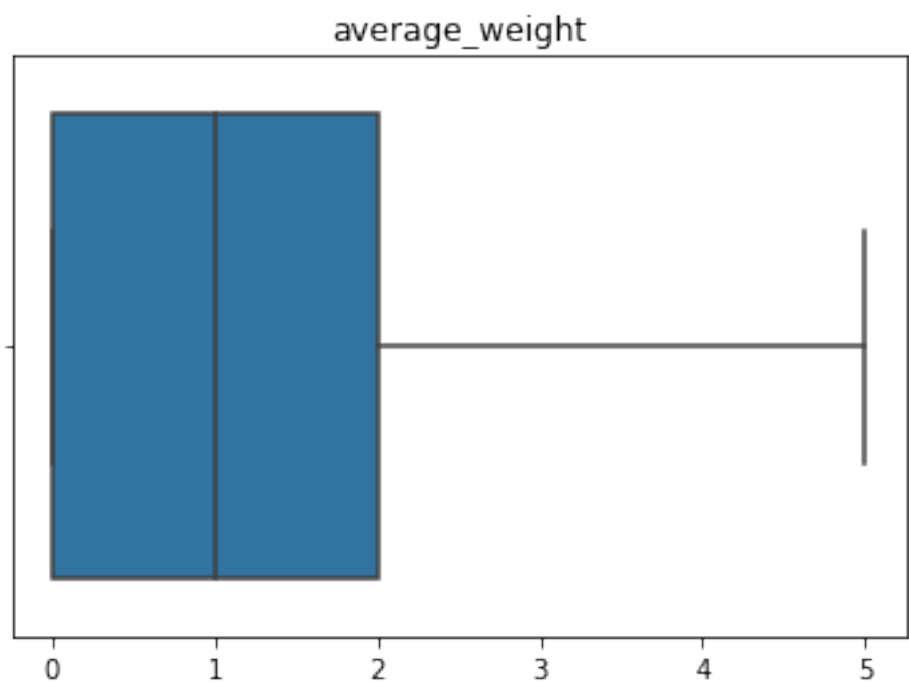













[39]: *#Input all the features into single vectors*

```
#standard Scaling all the feature vector to handle the outlier observed in the
→Boxplots
columnsReg=[c for c in df.columns if c not in
→['average_rating','yearpublished']]
assembler = VectorAssembler(inputCols=columnsReg, outputCol = 'Features')
output = assembler.transform(df)
scaled = StandardScaler(inputCol = 'Features',outputCol =
→'scaled_features',withStd = True,withMean = True)
final_df = scaled.fit(output).transform(output)
final_df.select('Features','scaled_features').show(1)
finalized_data = final_df.select("Features","average_rating","scaled_features")
finalized_data.show()
```

```
+-----+-----+
|          Features|      scaled_features|
+-----+-----+
|[2.0,2.0,180.0,18...|[-0.0063462794810...|
+-----+-----+
```

only showing top 1 row

```
+-----+-----+-----+
|          Features|average_rating|      scaled_features|
+-----+-----+-----+
|[2.0,2.0,180.0,18...|      8.33774|[-0.0063462794810...|
|[2.0,5.0,150.0,60...|      8.28798|[-0.0063462794810...|
|[1.0,7.0,210.0,30...|      8.28994|[-1.2694999838810...|
|[2.0,4.0,240.0,24...|      8.20407|[-0.0063462794810...|
|[2.0,5.0,150.0,90...|      8.14261|[-0.0063462794810...|
|[1.0,5.0,150.0,30...|      8.11957|[-1.2694999838810...|
|[2.0,2.0,45.0,45...|      8.1676|[-0.0063462794810...|
|[1.0,4.0,150.0,15...|      8.15901|[-1.2694999838810...|
|[2.0,4.0,90.0,30...|      8.07879|[-0.0063462794810...|
|[2.0,6.0,200.0,60...|      8.07933|[-0.0063462794810...|
|[2.0,6.0,120.0,12...|       7.9888|[-0.0063462794810...|
|[2.0,5.0,90.0,90...|      8.43944|[-0.0063462794810...|
|[2.0,4.0,150.0,15...|      8.35044|[-0.0063462794810...|
|[1.0,4.0,180.0,90...|      8.09283|[-1.2694999838810...|
|[1.0,5.0,200.0,10...|      7.99115|[-1.2694999838810...|
|[3.0,4.0,180.0,12...|      8.03071|[1.25680742491897...|
|[2.0,4.0,90.0,90...|      7.98673|[-0.0063462794810...|
|[2.0,5.0,210.0,45...|      8.05776|[-0.0063462794810...|
|[2.0,7.0,30.0,30...|      7.87047|[-0.0063462794810...|
|[2.0,5.0,150.0,60...|      7.89829|[-0.0063462794810...|
+-----+-----+-----+
```

only showing top 20 rows

```
[40]: #test Train split the Data into 80:20 Ratio
train,test = finalized_data.randomSplit([0.8,0.2],seed=42)
regressor = LinearRegression(featuresCol = 'scaled_features', labelCol = 'average_rating')

#Fitting the Training sent to linearRegression Model
regressor = regressor.fit(train)

#Model Building
pred = regressor.evaluate(test)

#Predicting
pred.predictions.show()
```

Features	average_rating	scaled_features	prediction
(14, [0,1,2,3,4,5,...]	3.0	[-1.2694999838810...	5.357054248980058
(14, [0,1,2,3,4,5,...]	2.0	[-1.2694999838810...	5.472450147939093
(14, [0,1,2,3,4,5,...]	4.0	[-1.2694999838810...	5.472450147939093
(14, [0,1,2,3,4,5,...]	7.0	[-0.0063462794810...	5.395723754488879
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.434594767362343
(14, [0,1,2,3,4,5,...]	2.0	[-0.0063462794810...	5.39565472239216
(14, [0,1,2,3,4,5,...]	4.0	[-0.0063462794810...	5.434334384004851
(14, [0,1,2,3,4,5,...]	1.0	[-0.0063462794810...	5.511311004708688
(14, [0,1,2,3,4,5,...]	5.0	[-0.0063462794810...	5.511311004708688
(14, [0,1,2,3,4,5,...]	1.0	[-0.0063462794810...	5.39622877617448
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.434717086526398
(14, [0,1,2,3,4,5,...]	5.0	[-0.0063462794810...	5.512267761012554
(14, [0,1,2,3,4,5,...]	6.0	[-0.0063462794810...	5.356844869100723
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.3959072332349605
(14, [0,1,2,3,4,5,...]	5.0	[-0.0063462794810...	5.511433323872742
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.513469155644529
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.513469155644529
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.513469155644529
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.513469155644529
(14, [0,1,2,3,4,5,...]	3.0	[-0.0063462794810...	5.513469155644529

only showing top 20 rows

```
[41]: #calculating the summary to find out the Root mean squared error and Rsquared to
      ↪check the model accuracy
Summary = regressor.summary
print("RMSE: %f" % Summary.rootMeanSquaredError)
print("r2: %f" % Summary.r2)
```



RMSE: 1.452548

r2: 0.155437

```
[45]: !sudo apt-get install texlive-xetex texlive-fonts-recommended
      ↪texlive-generic-recommended

      !jupyter nbconvert --to pdf
```

Reading package lists... Done

Building dependency tree

Reading state information... Done

The following additional packages will be installed:

fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono fonts-texgyre  
javascript-common libcupsfilters1 libcupsimage2 libgs9 libgs9-common  
libijs-0.35 libjbig2dec0 libjs-jquery libkpathsea6 libpotrace0 libptexenc1  
libruby2.5 libsynchronet1 libtexlua52 libtexlua52 libzzip-0-13 lmodern  
poppler-data preview-latex-style rake ruby ruby-did-you-mean ruby-minitest  
ruby-net-telnet ruby-power-assert ruby-test-unit ruby2.5  
rubygems-integration tlutils tex-common tex-gyre texlive-base  
texlive-binaries texlive-latex-base texlive-latex-extra  
texlive-latex-recommended texlive-pictures texlive-plain-generic tipa

Suggested packages:

fonts-noto apache2 | lighttpd | httpd poppler-utils ghostscript  
fonts-japanese-mincho | fonts-ipafont-mincho fonts-japanese-gothic  
| fonts-ipafont-gothic fonts-arphic-ukai fonts-arphic-uming fonts-nanum ri  
ruby-dev bundler debhelper gv | postscript-viewer perl-tk xpdf-reader  
| pdf-viewer texlive-fonts-recommended-doc texlive-latex-base-doc  
python-pygments icc-profiles libfile-which-perl  
libspreadsheet-parseexcel-perl texlive-latex-extra-doc  
texlive-latex-recommended-doc texlive-pstricks dot2tex prerex ruby-tcltk  
| libtcltk-ruby texlive-pictures-doc vprerex

The following NEW packages will be installed:

fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono fonts-texgyre  
javascript-common libcupsfilters1 libcupsimage2 libgs9 libgs9-common  
libijs-0.35 libjbig2dec0 libjs-jquery libkpathsea6 libpotrace0 libptexenc1  
libruby2.5 libsynchronet1 libtexlua52 libtexlua52 libzzip-0-13 lmodern  
poppler-data preview-latex-style rake ruby ruby-did-you-mean ruby-minitest  
ruby-net-telnet ruby-power-assert ruby-test-unit ruby2.5  
rubygems-integration tlutils tex-common tex-gyre texlive-base  
texlive-binaries texlive-fonts-recommended texlive-generic-recommended  
texlive-latex-base texlive-latex-extra texlive-latex-recommended  
texlive-pictures texlive-plain-generic texlive-xetex tipa

0 upgraded, 47 newly installed, 0 to remove and 14 not upgraded.

Need to get 146 MB of archives.

After this operation, 460 MB of additional disk space will be used.

Get:1 http://archive.ubuntu.com/ubuntu bionic/main amd64 fonts-droid-fallback  
all 1:6.0.1r16-1.1 [1,805 kB]

Get:2 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 fonts-lato all 2.0-2  
[2,698 kB]  
Get:3 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 poppler-data all  
0.4.8-2 [1,479 kB]  
Get:4 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 tex-common all 6.09  
[33.0 kB]  
Get:5 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 fonts-lmodern all  
2.004.5-3 [4,551 kB]  
Get:6 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 fonts-noto-mono all  
20171026-2 [75.5 kB]  
Get:7 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 fonts-texgyre all  
20160520-1 [8,761 kB]  
Get:8 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 javascript-common all  
11 [6,066 B]  
Get:9 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libcupsfilters1  
amd64 1.20.2-0ubuntu3.1 [108 kB]  
Get:10 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libcupsimage2  
amd64 2.2.7-1ubuntu2.8 [18.6 kB]  
Get:11 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 libijs-0.35 amd64  
0.35-13 [15.5 kB]  
Get:12 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 libjbig2dec0 amd64  
0.13-6 [55.9 kB]  
Get:13 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libgs9-common  
all 9.26~dfsg+0-0ubuntu0.18.04.13 [5,092 kB]  
Get:14 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libgs9 amd64  
9.26~dfsg+0-0ubuntu0.18.04.13 [2,263 kB]  
Get:15 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 libjs-jquery all  
3.2.1-1 [152 kB]  
Get:16 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libkpathsea6  
amd64 2017.20170613.44572-8ubuntu0.1 [54.9 kB]  
Get:17 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 libpotrace0 amd64  
1.14-2 [17.4 kB]  
Get:18 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libptexenc1  
amd64 2017.20170613.44572-8ubuntu0.1 [34.5 kB]  
Get:19 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 rubygems-integration  
all 1.11 [4,994 B]  
Get:20 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 ruby2.5 amd64  
2.5.1-1ubuntu1.7 [48.6 kB]  
Get:21 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby amd64 1:2.5.1  
[5,712 B]  
Get:22 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 rake all  
12.3.1-1ubuntu0.1 [44.9 kB]  
Get:23 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby-did-you-mean all  
1.2.0-2 [9,700 B]  
Get:24 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby-minitest all  
5.10.3-1 [38.6 kB]  
Get:25 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby-net-telnet all  
0.1.1-2 [12.6 kB]

Get:26 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby-power-assert all 0.3.0-1 [7,952 B]  
Get:27 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 ruby-test-unit all 3.2.5-1 [61.1 kB]  
Get:28 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libruby2.5 amd64 2.5.1-1ubuntu1.7 [3,068 kB]  
Get:29 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libsyntax1 amd64 2017.20170613.44572-8ubuntu0.1 [41.4 kB]  
Get:30 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libtexlua52 amd64 2017.20170613.44572-8ubuntu0.1 [91.2 kB]  
Get:31 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libtexluajit2 amd64 2017.20170613.44572-8ubuntu0.1 [230 kB]  
Get:32 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 libzip-0-13 amd64 0.13.62-3.1ubuntu0.18.04.1 [26.0 kB]  
Get:33 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 lmodern all 2.004.5-3 [9,631 kB]  
Get:34 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 preview-latex-style all 11.91-1ubuntu1 [185 kB]  
Get:35 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 t1utils amd64 1.41-2 [56.0 kB]  
Get:36 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 tex-gyre all 20160520-1 [4,998 kB]  
Get:37 <http://archive.ubuntu.com/ubuntu> bionic-updates/main amd64 texlive-binaries amd64 2017.20170613.44572-8ubuntu0.1 [8,179 kB]  
Get:38 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 texlive-base all 2017.20180305-1 [18.7 MB]  
Get:39 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-fonts-recommended all 2017.20180305-1 [5,262 kB]  
Get:40 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-generic-generic all 2017.20180305-2 [23.6 MB]  
Get:41 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-generic-recommended all 2017.20180305-1 [15.9 kB]  
Get:42 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 texlive-latex-base all 2017.20180305-1 [951 kB]  
Get:43 <http://archive.ubuntu.com/ubuntu> bionic/main amd64 texlive-latex-recommended all 2017.20180305-1 [14.9 MB]  
Get:44 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-pictures all 2017.20180305-1 [4,026 kB]  
Get:45 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-latex-extra all 2017.20180305-2 [10.6 MB]  
Get:46 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 tipa all 2:1.3-20 [2,978 kB]  
Get:47 <http://archive.ubuntu.com/ubuntu> bionic/universe amd64 texlive-xetex all 2017.20180305-1 [10.7 MB]  
Fetched 146 MB in 8s (17.5 MB/s)  
debconf: unable to initialize frontend: Dialog  
debconf: (No usable dialog-like program is installed, so the dialog based frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line 76,

```

<> line 47.)
debconf: falling back to frontend: Readline
debconf: unable to initialize frontend: Readline
debconf: (This frontend requires a controlling tty.)
debconf: falling back to frontend: Teletype
dpkg-preconfigure: unable to re-open stdin:
Selecting previously unselected package fonts-droid-fallback.
(Reading database ... 145208 files and directories currently installed.)
Preparing to unpack .../00-fonts-droid-fallback_1%3a6.0.1r16-1.1_all.deb ...
Unpacking fonts-droid-fallback (1:6.0.1r16-1.1) ...
Selecting previously unselected package fonts-lato.
Preparing to unpack .../01-fonts-lato_2.0-2_all.deb ...
Unpacking fonts-lato (2.0-2) ...
Selecting previously unselected package poppler-data.
Preparing to unpack .../02-poppler-data_0.4.8-2_all.deb ...
Unpacking poppler-data (0.4.8-2) ...
Selecting previously unselected package tex-common.
Preparing to unpack .../03-tex-common_6.09_all.deb ...
Unpacking tex-common (6.09) ...
Selecting previously unselected package fonts-lmodern.
Preparing to unpack .../04-fonts-lmodern_2.004.5-3_all.deb ...
Unpacking fonts-lmodern (2.004.5-3) ...
Selecting previously unselected package fonts-noto-mono.
Preparing to unpack .../05-fonts-noto-mono_20171026-2_all.deb ...
Unpacking fonts-noto-mono (20171026-2) ...
Selecting previously unselected package fonts-texgyre.
Preparing to unpack .../06-fonts-texgyre_20160520-1_all.deb ...
Unpacking fonts-texgyre (20160520-1) ...
Selecting previously unselected package javascript-common.
Preparing to unpack .../07-javascript-common_11_all.deb ...
Unpacking javascript-common (11) ...
Selecting previously unselected package libcupsfilters1:amd64.
Preparing to unpack .../08-libcupsfilters1_1.20.2-0ubuntu3.1_amd64.deb ...
Unpacking libcupsfilters1:amd64 (1.20.2-0ubuntu3.1) ...
Selecting previously unselected package libcupsimage2:amd64.
Preparing to unpack .../09-libcupsimage2_2.2.7-1ubuntu2.8_amd64.deb ...
Unpacking libcupsimage2:amd64 (2.2.7-1ubuntu2.8) ...
Selecting previously unselected package libijs-0.35:amd64.
Preparing to unpack .../10-libijs-0.35_0.35-13_amd64.deb ...
Unpacking libijs-0.35:amd64 (0.35-13) ...
Selecting previously unselected package libjbig2dec0:amd64.
Preparing to unpack .../11-libjbig2dec0_0.13-6_amd64.deb ...
Unpacking libjbig2dec0:amd64 (0.13-6) ...
Selecting previously unselected package libgs9-common.
Preparing to unpack .../12-libgs9-common_9.26~dfsg+0-0ubuntu0.18.04.13_all.deb
...
Unpacking libgs9-common (9.26~dfsg+0-0ubuntu0.18.04.13) ...
Selecting previously unselected package libgs9:amd64.

```

```

Preparing to unpack .../13-libgs9_9.26~dfsg+0-0ubuntu0.18.04.13_amd64.deb ...
Unpacking libgs:amd64 (9.26~dfsg+0-0ubuntu0.18.04.13) ...
Selecting previously unselected package libjs-jquery.
Preparing to unpack .../14-libjs-jquery_3.2.1-1_all.deb ...
Unpacking libjs-jquery (3.2.1-1) ...
Selecting previously unselected package libkpathsea6:amd64.
Preparing to unpack .../15-libkpathsea6_2017.20170613.44572-8ubuntu0.1_amd64.deb
...
Unpacking libkpathsea6:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package libpotrace0.
Preparing to unpack .../16-libpotrace0_1.14-2_amd64.deb ...
Unpacking libpotrace0 (1.14-2) ...
Selecting previously unselected package libptexenc1:amd64.
Preparing to unpack .../17-libptexenc1_2017.20170613.44572-8ubuntu0.1_amd64.deb
...
Unpacking libptexenc1:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package rubygems-integration.
Preparing to unpack .../18-rubygems-integration_1.11_all.deb ...
Unpacking rubygems-integration (1.11) ...
Selecting previously unselected package ruby2.5.
Preparing to unpack .../19-ruby2.5_2.5.1-1ubuntu1.7_amd64.deb ...
Unpacking ruby2.5 (2.5.1-1ubuntu1.7) ...
Selecting previously unselected package ruby.
Preparing to unpack .../20-ruby_1%3a2.5.1_amd64.deb ...
Unpacking ruby (1:2.5.1) ...
Selecting previously unselected package rake.
Preparing to unpack .../21-rake_12.3.1-1ubuntu0.1_all.deb ...
Unpacking rake (12.3.1-1ubuntu0.1) ...
Selecting previously unselected package ruby-did-you-mean.
Preparing to unpack .../22-ruby-did-you-mean_1.2.0-2_all.deb ...
Unpacking ruby-did-you-mean (1.2.0-2) ...
Selecting previously unselected package ruby-minitest.
Preparing to unpack .../23-ruby-minitest_5.10.3-1_all.deb ...
Unpacking ruby-minitest (5.10.3-1) ...
Selecting previously unselected package ruby-net-telnet.
Preparing to unpack .../24-ruby-net-telnet_0.1.1-2_all.deb ...
Unpacking ruby-net-telnet (0.1.1-2) ...
Selecting previously unselected package ruby-power-assert.
Preparing to unpack .../25-ruby-power-assert_0.3.0-1_all.deb ...
Unpacking ruby-power-assert (0.3.0-1) ...
Selecting previously unselected package ruby-test-unit.
Preparing to unpack .../26-ruby-test-unit_3.2.5-1_all.deb ...
Unpacking ruby-test-unit (3.2.5-1) ...
Selecting previously unselected package libruby2.5:amd64.
Preparing to unpack .../27-libruby2.5_2.5.1-1ubuntu1.7_amd64.deb ...
Unpacking libruby2.5:amd64 (2.5.1-1ubuntu1.7) ...
Selecting previously unselected package libsyntax1:amd64.
Preparing to unpack .../28-libsyntax1_2017.20170613.44572-8ubuntu0.1_amd64.deb

```

```

...
Unpacking libsynctex1:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package libtexlua52:amd64.
Preparing to unpack .../29-libtexlua52_2017.20170613.44572-8ubuntu0.1_amd64.deb
...
Unpacking libtexlua52:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package libtexluajit2:amd64.
Preparing to unpack
.../30-libtexluajit2_2017.20170613.44572-8ubuntu0.1_amd64.deb ...
Unpacking libtexluajit2:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package libzip-0-13:amd64.
Preparing to unpack .../31-libzip-0-13_0.13.62-3.1ubuntu0.18.04.1_amd64.deb ...
Unpacking libzip-0-13:amd64 (0.13.62-3.1ubuntu0.18.04.1) ...
Selecting previously unselected package lmodern.
Preparing to unpack .../32-lmodern_2.004.5-3_all.deb ...
Unpacking lmodern (2.004.5-3) ...
Selecting previously unselected package preview-latex-style.
Preparing to unpack .../33-preview-latex-style_11.91-1ubuntu1_all.deb ...
Unpacking preview-latex-style (11.91-1ubuntu1) ...
Selecting previously unselected package tlutils.
Preparing to unpack .../34-tlutils_1.41-2_amd64.deb ...
Unpacking tlutils (1.41-2) ...
Selecting previously unselected package tex-gyre.
Preparing to unpack .../35-tex-gyre_20160520-1_all.deb ...
Unpacking tex-gyre (20160520-1) ...
Selecting previously unselected package texlive-binaries.
Preparing to unpack .../36-texlive-
binaries_2017.20170613.44572-8ubuntu0.1_amd64.deb ...
Unpacking texlive-binaries (2017.20170613.44572-8ubuntu0.1) ...
Selecting previously unselected package texlive-base.
Preparing to unpack .../37-texlive-base_2017.20180305-1_all.deb ...
Unpacking texlive-base (2017.20180305-1) ...
Selecting previously unselected package texlive-fonts-recommended.
Preparing to unpack .../38-texlive-fonts-recommended_2017.20180305-1_all.deb ...
Unpacking texlive-fonts-recommended (2017.20180305-1) ...
Selecting previously unselected package texlive-plain-generic.
Preparing to unpack .../39-texlive-plain-generic_2017.20180305-2_all.deb ...
Unpacking texlive-plain-generic (2017.20180305-2) ...
Selecting previously unselected package texlive-generic-recommended.
Preparing to unpack .../40-texlive-generic-recommended_2017.20180305-1_all.deb
...
Unpacking texlive-generic-recommended (2017.20180305-1) ...
Selecting previously unselected package texlive-latex-base.
Preparing to unpack .../41-texlive-latex-base_2017.20180305-1_all.deb ...
Unpacking texlive-latex-base (2017.20180305-1) ...
Selecting previously unselected package texlive-latex-recommended.
Preparing to unpack .../42-texlive-latex-recommended_2017.20180305-1_all.deb ...
Unpacking texlive-latex-recommended (2017.20180305-1) ...

```

```

Selecting previously unselected package texlive-pictures.
Preparing to unpack .../43-texlive-pictures_2017.20180305-1_all.deb ...
Unpacking texlive-pictures (2017.20180305-1) ...
Selecting previously unselected package texlive-latex-extra.
Preparing to unpack .../44-texlive-latex-extra_2017.20180305-2_all.deb ...
Unpacking texlive-latex-extra (2017.20180305-2) ...
Selecting previously unselected package tipa.
Preparing to unpack .../45-tipa_2%3a1.3-20_all.deb ...
Unpacking tipa (2:1.3-20) ...
Selecting previously unselected package texlive-xetex.
Preparing to unpack .../46-texlive-xetex_2017.20180305-1_all.deb ...
Unpacking texlive-xetex (2017.20180305-1) ...
Setting up libgs9-common (9.26~dfsg+0-0ubuntu0.18.04.13) ...
Setting up libkpathsea6:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Setting up libjs-jquery (3.2.1-1) ...
Setting up libtexlua52:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Setting up fonts-droid-fallback (1:6.0.1r16-1.1) ...
Setting up libsynchronet1:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Setting up libptexenc1:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Setting up tex-common (6.09) ...
debconf: unable to initialize frontend: Dialog
debconf: (No usable dialog-like program is installed, so the dialog based
frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line
76.)
debconf: falling back to frontend: Readline
update-language: texlive-base not installed and configured, doing nothing!
Setting up poppler-data (0.4.8-2) ...
Setting up tex-gyre (20160520-1) ...
Setting up preview-latex-style (11.91-1ubuntu1) ...
Setting up fonts-texgyre (20160520-1) ...
Setting up fonts- noto-mono (20171026-2) ...
Setting up fonts-lato (2.0-2) ...
Setting up libcupsfilters1:amd64 (1.20.2-0ubuntu3.1) ...
Setting up libcupsimage2:amd64 (2.2.7-1ubuntu2.8) ...
Setting up libjbig2dec0:amd64 (0.13-6) ...
Setting up ruby-did-you-mean (1.2.0-2) ...
Setting up t1utils (1.41-2) ...
Setting up ruby-net-telnet (0.1.1-2) ...
Setting up libijs-0.35:amd64 (0.35-13) ...
Setting up rubygems-integration (1.11) ...
Setting up libpotrace0 (1.14-2) ...
Setting up javascript-common (11) ...
Setting up ruby-minitest (5.10.3-1) ...
Setting up libzzip-0-13:amd64 (0.13.62-3.1ubuntu0.18.04.1) ...
Setting up libgs9:amd64 (9.26~dfsg+0-0ubuntu0.18.04.13) ...
Setting up libtexluajit2:amd64 (2017.20170613.44572-8ubuntu0.1) ...
Setting up fonts-lmodern (2.004.5-3) ...
Setting up ruby-power-assert (0.3.0-1) ...

```

```

Setting up texlive-binaries (2017.20170613.44572-8ubuntu0.1) ...
update-alternatives: using /usr/bin/xdvi-xaw to provide /usr/bin/xdvi.bin
(xdvi.bin) in auto mode
update-alternatives: using /usr/bin/bibtex.original to provide /usr/bin/bibtex
(bibtex) in auto mode
Setting up texlive-base (2017.20180305-1) ...
mktexlsr: Updating /var/lib/texmf/ls-R-TEXLIVEDIST...
mktexlsr: Updating /var/lib/texmf/ls-R-TEXMFMAIN...
mktexlsr: Updating /var/lib/texmf/ls-R...
mktexlsr: Done.
tl-paper: setting paper size for dvips to a4: /var/lib/texmf/dvips/config
/config-paper.ps
tl-paper: setting paper size for dvipdfmx to a4: /var/lib/texmf/dvipdfmx
/dvipdfmx-paper.cfg
tl-paper: setting paper size for xdvi to a4: /var/lib/texmf/xdvi/XDvi-paper
tl-paper: setting paper size for pdftex to a4:
/var/lib/texmf/tex/generic/config/pdftexconfig.tex
debconf: unable to initialize frontend: Dialog
debconf: (No usable dialog-like program is installed, so the dialog based
frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line
76.)
debconf: falling back to frontend: Readline
Setting up texlive-fonts-recommended (2017.20180305-1) ...
Setting up texlive-plain-generic (2017.20180305-2) ...
Setting up texlive-generic-recommended (2017.20180305-1) ...
Setting up texlive-latex-base (2017.20180305-1) ...
Setting up lmodern (2.004.5-3) ...
Setting up texlive-latex-recommended (2017.20180305-1) ...
Setting up texlive-pictures (2017.20180305-1) ...
Setting up tipa (2:1.3-20) ...
Regenerating '/var/lib/texmf/fmtutil.cnf-DEBIAN'... done.
Regenerating '/var/lib/texmf/fmtutil.cnf-TEXLIVEDIST'... done.
update-fmtutil has updated the following file(s):
    /var/lib/texmf/fmtutil.cnf-DEBIAN
    /var/lib/texmf/fmtutil.cnf-TEXLIVEDIST
If you want to activate the changes in the above file(s),
you should run fmtutil-sys or fmtutil.
Setting up texlive-latex-extra (2017.20180305-2) ...
Setting up texlive-xetex (2017.20180305-1) ...
Setting up ruby2.5 (2.5.1-1ubuntu1.7) ...
Setting up ruby (1:2.5.1) ...
Setting up ruby-test-unit (3.2.5-1) ...
Setting up rake (12.3.1-1ubuntu0.1) ...
Setting up libruby2.5:amd64 (2.5.1-1ubuntu1.7) ...
Processing triggers for mime-support (3.60ubuntu1) ...
Processing triggers for libc-bin (2.27-3ubuntu1.2) ...
/sbin/ldconfig.real: /usr/local/lib/python3.6/dist-
packages/ideep4py/lib/libmkldnn.so.0 is not a symbolic link

```



```

Processing triggers for man-db (2.8.3-2ubuntu0.1) ...
Processing triggers for fontconfig (2.12.6-0ubuntu2) ...
Processing triggers for tex-common (6.09) ...
debconf: unable to initialize frontend: Dialog
debconf: (No usable dialog-like program is installed, so the dialog based
frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line
76.)
debconf: falling back to frontend: Readline
Running updmap-sys. This may take some time... done.
Running mktexlsr /var/lib/texmf ... done.
Building format(s) --all.
    This may take some time... done.
This application is used to convert notebook files (*.ipynb) to various other
formats.

```

WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.

## Options

-----

Arguments that take values are actually convenience aliases to full  
Configurables, whose aliases are listed on the help line. For more information  
on full configurables, see '--help-all'.

### --execute

Execute the notebook prior to export.

### --allow-errors

Continue notebook execution even if one of the cells throws an error and  
include the error message in the cell output (the default behaviour is to abort  
conversion). This flag is only relevant if '--execute' was specified, too.

### --no-input

Exclude input cells and output prompts from converted document.

This mode is ideal for generating code-free reports.

### --stdout

Write notebook output to stdout instead of files.

### --stdin

read a single notebook file from stdin. Write the resulting notebook with  
default basename 'notebook.\*'

### --inplace

Run nbconvert in place, overwriting the existing notebook (only  
relevant when converting to notebook format)

### -y

Answer yes to any questions instead of prompting.

### --clear-output

Clear output of current file and save in place,  
overwriting the existing notebook.

### --debug

```

    set log level to logging.DEBUG (maximize logging output)
--no-prompt
    Exclude input and output prompts from converted document.
--generate-config
    generate default config file
--nbformat=<Enum> (NotebookExporter.nbformat_version)
    Default: 4
    Choices: [1, 2, 3, 4]
    The nbformat version to write. Use this to downgrade notebooks.
--output-dir=<Unicode> (FilesWriter.build_directory)
    Default: ''
    Directory to write output(s) to. Defaults to output to the directory of each
    notebook. To recover previous default behaviour (outputting to the current
    working directory) use . as the flag value.
--writer=<DottedObjectName> (NbConvertApp.writer_class)
    Default: 'FilesWriter'
    Writer class used to write the results of the conversion
--log-level=<Enum> (Application.log_level)
    Default: 30
    Choices: (0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR',
'CRITICAL')
    Set the log level by value or name.
--reveal-prefix=<Unicode> (SlidesExporter.reveal_url_prefix)
    Default: u''
    The URL prefix for reveal.js (version 3.x). This defaults to the reveal CDN,
    but can be any url pointing to a copy of reveal.js.
    For speaker notes to work, this must be a relative path to a local copy of
    reveal.js: e.g., "reveal.js".
    If a relative path is given, it must be a subdirectory of the current
    directory (from which the server is run).
    See the usage documentation
    (https://nbconvert.readthedocs.io/en/latest/usage.html#reveal-js-html-slideshow)
    for more details.
--to=<Unicode> (NbConvertApp.export_format)
    Default: 'html'
    The export format to be used, either one of the built-in formats
    ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf',
    'python', 'rst', 'script', 'slides'] or a dotted object name that represents
    the import path for an `Exporter` class
--template=<Unicode> (TemplateExporter.template_file)
    Default: u''
    Name of the template file to use
--output=<Unicode> (NbConvertApp.output_base)
    Default: ''
    overwrite base name use for output files. can only be used when converting
    one notebook at a time.
--post=<DottedOrNone> (NbConvertApp.postprocessor_class)
    Default: u''

```

PostProcessor class used to write the results of the conversion  
--config=<Unicode> (JupyterApp.config\_file)  
Default: u''  
Full path of a config file.

To see all available configurables, use `--help-all`

## Examples

-----

The simplest way to use nbconvert is

```
> jupyter nbconvert mynotebook.ipynb
```

which will convert mynotebook.ipynb to the default format (probably HTML).

You can specify the export format with `--to`.

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'].

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes 'base', 'article' and 'report'. HTML includes 'basic' and 'full'. You can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template basic mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb  
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

```
[47]: !jupyter nbconvert --to pdf RegressionSpark.ipynb
```

```
[NbConvertApp] WARNING | pattern u'RegressionSpark.ipynb' matched no files  
This application is used to convert notebook files (*.ipynb) to various other  
formats.
```

```
WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.
```

```
Options
```

```
-----
```

```
Arguments that take values are actually convenience aliases to full  
Configurables, whose aliases are listed on the help line. For more information  
on full configurables, see '--help-all'.
```

```
--execute
```

```
    Execute the notebook prior to export.
```

```
--allow-errors
```

```
    Continue notebook execution even if one of the cells throws an error and  
include the error message in the cell output (the default behaviour is to abort  
conversion). This flag is only relevant if '--execute' was specified, too.
```

```
--no-input
```

```
    Exclude input cells and output prompts from converted document.
```

```
    This mode is ideal for generating code-free reports.
```

```
--stdout
```

```
    Write notebook output to stdout instead of files.
```

```
--stdin
```

```
    read a single notebook file from stdin. Write the resulting notebook with  
default basename 'notebook.*'
```

```
--inplace
```

```
    Run nbconvert in place, overwriting the existing notebook (only  
relevant when converting to notebook format)
```

```
-y
```

```
    Answer yes to any questions instead of prompting.
```

```
--clear-output
```

```
    Clear output of current file and save in place,  
    overwriting the existing notebook.
```

```
--debug
```

```
    set log level to logging.DEBUG (maximize logging output)
```

```
--no-prompt
```

```
    Exclude input and output prompts from converted document.
```

```
--generate-config
```

```

    generate default config file
--nbformat=<Enum> (NotebookExporter.nbformat_version)
    Default: 4
    Choices: [1, 2, 3, 4]
    The nbformat version to write. Use this to downgrade notebooks.
--output-dir=<Unicode> (FilesWriter.build_directory)
    Default: ''
    Directory to write output(s) to. Defaults to output to the directory of each
    notebook. To recover previous default behaviour (outputting to the current
    working directory) use . as the flag value.
--writer=<DottedObjectName> (NbConvertApp.writer_class)
    Default: 'FilesWriter'
    Writer class used to write the results of the conversion
--log-level=<Enum> (Application.log_level)
    Default: 30
    Choices: (0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR',
'CRITICAL')
    Set the log level by value or name.
--reveal-prefix=<Unicode> (SlidesExporter.reveal_url_prefix)
    Default: u''
    The URL prefix for reveal.js (version 3.x). This defaults to the reveal CDN,
    but can be any url pointing to a copy of reveal.js.
    For speaker notes to work, this must be a relative path to a local copy of
    reveal.js: e.g., "reveal.js".
    If a relative path is given, it must be a subdirectory of the current
    directory (from which the server is run).
    See the usage documentation
    (https://nbconvert.readthedocs.io/en/latest/usage.html#reveal-js-html-slideshow) for more details.
--to=<Unicode> (NbConvertApp.export_format)
    Default: 'html'
    The export format to be used, either one of the built-in formats
    ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf',
    'python', 'rst', 'script', 'slides'] or a dotted object name that represents
    the import path for an `Exporter` class
--template=<Unicode> (TemplateExporter.template_file)
    Default: u''
    Name of the template file to use
--output=<Unicode> (NbConvertApp.output_base)
    Default: ''
    overwrite base name use for output files. can only be used when converting
    one notebook at a time.
--post=<DottedOrNone> (NbConvertApp.postprocessor_class)
    Default: u''
    PostProcessor class used to write the results of the conversion
--config=<Unicode> (JupyterApp.config_file)
    Default: u''
    Full path of a config file.

```

To see all available configurables, use `--help-all`

## Examples

-----

The simplest way to use nbconvert is

```
> jupyter nbconvert mynotebook.ipynb
```

which will convert mynotebook.ipynb to the default format (probably HTML).

You can specify the export format with `--to`.

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'].

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes 'base', 'article' and 'report'. HTML includes 'basic' and 'full'. You can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template basic mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

```
[49]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.